

Durability and Contention in Software Transactional Memory

DIPLOMARBEIT

zur Erlangung des akademischen Grades

Diplom-Ingenieur

im Rahmen des Studiums

Software Engineering & Internet Computing

eingereicht von

Michael Schröder

Matrikelnummer 0725699

an der Fakultät für Informatik
der Technischen Universität Wien

Betreuung: Univ. Prof. Dipl.-Inf. Dr. rer. nat. Jens Knoop
Mitwirkung: Assoc. Prof. Gabriele Keller, University of New South Wales

Wien, 1. August 2015

Michael Schröder

Jens Knoop

Durability and Contention in Software Transactional Memory

DIPLOMA THESIS

submitted in partial fulfillment of the requirements for the degree of

Master of Science

in

Software Engineering & Internet Computing

by

Michael Schröder

Registration Number 0725699

to the Faculty of Informatics
at the Vienna University of Technology

Advisor: Univ. Prof. Dipl.-Inf. Dr. rer. nat. Jens Knoop

Assistance: Assoc. Prof. Gabriele Keller, University of New South Wales

Vienna, 1st August, 2015

Michael Schröder

Jens Knoop

Erklärung zur Verfassung der Arbeit

Michael Schröder
Rosasgasse 15/8
1120 Wien

Hiermit erkläre ich, dass ich diese Arbeit selbständig verfasst habe, dass ich die verwendeten Quellen und Hilfsmittel vollständig angegeben habe und dass ich die Stellen der Arbeit – einschließlich Tabellen, Karten und Abbildungen –, die anderen Werken oder dem Internet im Wortlaut oder dem Sinn nach entnommen sind, auf jeden Fall unter Angabe der Quelle als Entlehnung kenntlich gemacht habe.

Wien, 1. August 2015

Michael Schröder

Acknowledgements

I thank Jens Knoop, my supervisor at the Vienna University of Technology, for his advice and support, as well as Gabriele Keller for inviting me to visit the University of New South Wales for three months to work on this thesis. I thank the Vienna University of Technology for supporting my research with a stipend. I also thank Manuel Chakravarty for suggesting the name “finalizers” and Simon Meier for pointing out some issues with the original transactional trie benchmarks.

Kurzfassung

Software Transactional Memory (STM) vereinfacht nebenläufige Programmierung, indem es erlaubt Speicheroperationen zu atomaren Blöcken zusammenzufassen. STM Transaktionen bieten Atomarität, Konsistenz und Isolation, und sind damit Datenbanktransaktionen ähnlich. Im Unterschied zu Datenbanken bietet STM allerdings keine *Dauerhaftigkeit*.

Ich habe Haskell's STM Implementierung mit einem Mechanismus erweitert, mit dem man *finalizer* zu atomaren Blöcken hinzufügen kann. Die neue Operation `atomicallyWithIO` erlaubt dem Programmierer die sichere Ausführung von beliebigen I/O-Operationen während der Commit-Phase einer Transaktion. Dies kann dazu benutzt werden STM dauerhaft zu machen, hat aber noch mehr Anwendungsmöglichkeiten. Zum Beispiel könnte ein *finalizer* den Benutzer fragen anstehende Resultate zu genehmigen, was interaktive Transaktionen ermöglicht.

Ein weiteres häufiges Problem mit STM ist *contention*. Viele Standarddatenstrukturen, wenn sie in einer transaktionalen Umgebung verwendet werden, erzeugen eine übermäßig hohe Anzahl an Konflikten. Ich stelle eine neue contention-freie STM-Datenstruktur vor, den *transactional trie*. Er basiert auf dem lock-free concurrent trie und benutzt lokale Seiteneffekte um unnötige Konflikte zu vermeiden, während er gleichzeitig die transaktionale Sicherheit bewahrt.

Sowohl *finalizer* als auch der *transactional trie* sind Beispiele für das Kombinieren von Transaktionen mit Seiteneffekten. *Finalizer* sind ein allgemeiner top-down Ansatz, während der *transactional trie* Seiteneffekte auf der Mikroebene integriert. Ich demonstriere die Effektivität von beidem in dem ich eine vollständige Beispielsanwendung baue die STM als Datenbanksprache benutzt, Dauerhaftigkeit bietet und contention vermeidet.

Abstract

Software Transactional Memory (STM) immensely simplifies concurrent programming by allowing memory operations to be grouped together into atomic blocks. Like database transactions, STM transactions provide atomicity, consistency and isolation. Unlike databases, they do not provide *durability*.

I extend Haskell's STM implementation with a mechanism to add *finalizers* to atomic blocks. The new operation `atomicallyWithIO` allows the programmer to safely execute arbitrary I/O during the commit-phase of a transaction. This can be used to make STM durable, but it turns out to have even more applications. For example, a finalizer could ask the user to approve pending results, enabling interactive transactions.

Another common problem with STM is *contention*. When used in a transactional setting, many standard data structures cause unreasonably high numbers of conflicts. I propose a new STM data structure, the *transactional trie*, a contention-free hash map. It is based on the lock-free concurrent trie, and uses localized side-effects to eliminate unnecessary conflicts while preserving transactional safety.

Both finalizers and the transactional trie are examples of combining transactions with side-effects. Finalizers are a general top-down approach, while the transactional trie incorporates side effects on the micro-level. I demonstrate the effectiveness of both by building a full sample application that uses STM as a database language, providing durability and avoiding contention.

Contents

1	Motivation	1
2	Transactional Memory with Finalizers	5
2.1	STM and ACID	5
2.2	Finalizers	7
2.3	Related Work	10
2.4	Semantics of STM	11
2.5	Semantics of finalizers	18
2.6	Implementation	27
3	Transactional Tries	31
3.1	Background	32
3.2	Implementation	34
3.3	Memory efficiency	39
3.4	Evaluation	42
4	STM as a database language	49
4.1	Example: a social network	49
4.2	The TX monad	52
4.3	Caveats	55
5	Conclusions & Perspectives	57
5.1	Related Work	57
5.2	Future Work	59
	Bibliography	61
	Publications	65

Motivation

It is a widely held opinion that concurrent programming is difficult and error-prone. Low-level synchronization mechanisms, such as locks, are notoriously tricky to get right. Deadlocks, livelocks, heisenbugs and other issues encountered when writing complex concurrent systems are usually hard to track down and often confound even experienced programmers.

To simplify concurrent programming, higher-level abstractions are needed. One such abstraction is *Software Transactional Memory* (STM). Briefly, this technique allows the programmer to group multiple memory operations into a single atomic block, not unlike a database transaction. When implemented in a high-level language such as Haskell, with its emphasis on purity and its strong static type system, STM becomes especially powerful.

However, it is not always as powerful as we would like it to be:

- While manipulating memory using STM is easy, persisting those manipulations in a transactionally safe way is frustratingly impossible. There is no way to achieve *durability*, an important component of database transactions.
- Some pure data types, such as maps, can be highly inefficient when combined with STM. The reason is *contention*: certain common access patterns cause unreasonably high numbers of transactional conflicts.

In this thesis, I will address these problems by combining transactions with side-effects. In particular, the contributions of my work are:

- A new STM primitive called `atomicallyWithIO` that allows the user to attach a *finalizer* to an STM transaction. Such finalizers are arbitrary I/O actions that can be used to serialize transactional data or incorporate external information into a transaction. I provide a detailed discussion of the semantics of this extension to STM, as well as my implementation of it in the Glasgow Haskell Compiler (GHC). (Chapter 2)

- A contention-free STM data structure, the *transactional trie*. It is based on the concurrent trie, but lifted into an STM context and combines transactions with carefully considered internal side effects. I present its design and implementation in Haskell, and evaluate it against other STM-specialized data structures. (Chapter 3)
- A demonstration of the effectiveness of both finalizers and transactional tries by constructing a thin database abstraction on top of STM, which I then use to build a simple social networking site. (Chapter 4)

In the remainder of this introductory chapter, I will give a brief overview of Haskell's STM interface.

Background: STM in Haskell

Here are the main data types and operations of STM in Haskell:

```
data STM a
instance Monad STM
atomically :: STM a → IO a
data TVar a
newTVar :: a → STM (TVar a)
readTVar :: TVar a → STM a
writeTVar :: TVar a → a → STM ()
retry :: STM a
orElse :: STM a → STM a → STM a
```

Atomic blocks in Haskell are represented by the STM monad. Inside this monad, we can freely operate on transactional variables, or **TVars**. We can read them, write them and create new ones. When we want to actually perform an STM computation and make its effects visible to the rest of the world, we apply **atomically** to the computation. This function turns an STM block into a transaction in the IO monad that, when executed, will take place atomically with respect to all other transactions.

For example, the following code snippet increments a transactional variable *v*:

```
atomically $ do x ← readTVar v
              writeTVar v (x + 1)
```

The use of **atomically** guarantees that no other thread can come in between the reading and writing of the variable. The sequence of operations happens indivisibly.

An important aspect of Haskell's STM implementation is that it is fully composable. Smaller transactions can be combined into larger transactions without having to know how these smaller transactions are implemented. An important tool to make this possible is the composable blocking operator **retry**. Conceptually, **retry** abandons the current transaction and runs it again from the top. In the following example, the variable *v* is

decremented, unless it is zero, in which case the transaction blocks until v is non-zero again:

```
atomically $ do  $x \leftarrow \text{readTVar } v$ 
             if  $x \equiv 0$ 
             then retry
             else  $\text{writeTVar } v (x - 1)$ 
```

In addition to `retry`, there is the `orElse` combinator, which allows “trying out” transactions in sequence. `m1 'orElse' m2` first executes `m1`; if `m1` returns, then `orElse` returns; but if `m1` retries, its effects are discarded and `m2` is executed instead.

STM is also robust against exceptions. The standard functions `throw` and `catch` act as expected: if an exception occurs inside an atomic block and is not caught, the transaction’s effects are discarded and the exception is propagated.

GHC’s implementation of STM also contains support for data invariants. Using the `alwaysSucceeds` function, one can introduce an invariant over transactional variables that is dynamically checked on every atomic update.

For more background on Haskell’s STM, including its implementation, see the original STM papers (Harris, Marlow, et al. 2005; Harris and Peyton Jones 2006). For a more thorough exploration of not only STM but also other Haskell concurrency mechanisms, read Simon Marlow’s excellent book on that topic (Marlow 2013b).

Transactional Memory with Finalizers

2.1 STM and ACID

Atomicity, consistency, isolation and durability (ACID) are the cornerstones of any database system. Haskell’s STM gives us three of these properties: transactions are always *atomic* and have an *isolated* view of memory; dynamically-checked invariants can be used to ensure that the state of the system is *consistent*. The one missing property is *durability*.

Making data durable means serializing it to storage. This is of course an I/O action, which is not allowed within the confines of STM. For good reason: STM transactions can retry at any time and I/O actions are in general neither idempotent nor interruptible. However, notice that for the purposes of durability it is not necessary to run arbitrary I/O at arbitrary moments during a transaction. We merely need to run a single I/O action to serialize our data at the end of the transaction.

The naive approach is to first atomically perform some STM computation and then independently serialize its result:

```
durably :: STM a → IO ()
durably m = do x ← atomically m
             serialize x
```

There are two issues with this. First, by virtue of the serialization happening independently of the atomic block, after we have performed a transaction another thread could perform a second transaction and serialize it before we have finished serializing the first one. Depending on our serialization method, we could end up with an inconsistent state between the memory of our program and what is stored on disk. At the very least, there is no guarantee that the ordering of events is preserved. Secondly, the function `serialize`

might not terminate at all; it could throw an exception or its thread could die. Again we would end up with an inconsistent state and possibly data loss.

So the serialization needs to be somehow coupled with the atomic transaction in such a way that the transaction only commits *after* the data has been serialized. We might try to use the ominously named `unsafeIOToSTM` function:

```
durably :: STM a → IO ()
durably m = atomically $ do x ← m
                        unsafeIOToSTM (serialize x)
```

The type of `unsafeIOToSTM` is `IO a → STM a`, effectively allowing us to circumvent the type system so we can unsafely perform I/O in the STM monad. But this too must fail: even if it is the last statement in the atomic block, `serialize` is performed before the transaction commits. The transaction might yet have to retry because of a conflict with another transaction. In doing so it also executes `serialize` again. Unless `serialize` is idempotent, this is certainly undesirable. Furthermore, if the thread receives an asynchronous exception, the transaction will abort in an orderly fashion, while `serialize`, with its irrevocable side effects, cannot be undone.

This leads us to the realization that *the transaction must only commit after the data has been serialized, and the data must only be serialized after the transaction has committed*. We can escape the paradox by clarifying the second restriction: the data can in fact already be serialized after the transaction is merely *guaranteed* to commit. Meaning: at the point of serialization, the transaction has not yet made its effects visible to other transactions, but there are no conflicts preventing it from doing so. And until serialization is finished there must be no way for any such conflicts to arise; for all intents and purposes, the transaction must count as committed, even though its effects are not visible yet and it could still roll back by its own volition (but not due to conflicts with other transactions).

What is needed, then, is a new STM primitive. Not a combination of existing functions, but a new fundamental operation that does exactly what we want. It would have to be implemented directly in the runtime system and we would have to make sure that it is sound and does not allow us to undermine the safety of STM in general.

I will now describe this primitive, give its detailed semantics and talk about my implementation of it in the Glasgow Haskell Compiler.

2.2 Finalizers

The main idea is to introduce a new function

`atomicallyWithIO :: STM a → (a → IO b) → IO b`

Like the existing `atomically` operation, `atomicallyWithIO` transforms a computation of type `STM a` into an I/O action. Additionally, it takes a *finalizer*, which is a function of type `a → IO b`, viz. an I/O action that can depend on the result of the STM computation, and combines it with the transaction in such a way that:

1. The finalizer is only performed if the STM transaction is guaranteed to commit.
2. The STM transaction only commits (i.e. makes its effects visible to other transactions) if the finalizer finishes without raising an exception.

A detailed specification of `atomicallyWithIO` will be given in Section 2.5. Everything described there has been fully implemented by me in GHC and I will discuss that implementation in Section 2.6. But first let us look at some motivating examples, which will demonstrate the usefulness of this new operation and highlight some of its finer issues.

2.2.1 Example 1: Printing tickets

Say we want to sell tickets for an event. We have a set number of tickets, stored in a transactional variable `tickets` of type `TVar Int`, and the following function to get the next available ticket:

```
nextTicket :: STM Int
nextTicket = do t ← readTVar tickets
              when (t ≡ 0) (error "sold out")
              writeTVar tickets (t - 1)
              return t
```

Our box offices can then use a statement like `t ← atomically nextTicket` to safely acquire a ticket.

However, selling tickets involves actually printing the ticket number on a piece of paper. What happens if the printer malfunctions, or runs out of paper? If for some reason the ticket cannot be printed, we do not want that ticket number to be lost forever; it has not yet been sold after all. The solution, of course, is to use a finalizer:

```
printTicket :: Int → IO ()
printTicket t = ...
sellTicket :: IO ()
sellTicket = atomicallyWithIO nextTicket printTicket
```

If `printTicket` throws an exception, the transaction rolls back and the `tickets` counter was never decreased.

2.2.2 Example 2: User input

An I/O action can of course collect input as well. Consider the following example of an ATM. The function to withdraw money from an account is quite standard; as is, one would imagine, the function to dispense actual cash from the machine:

```
type Account = TVar Int
withdraw :: Account → Int → STM ()
withdraw acc amount = do bal ← readTVar acc
                      if bal < amount
                      then error "insufficient funds"
                      else writeTVar acc (bal - amount)

dispenseCash :: Int → IO ()
dispenseCash amount = ...
```

But instead of simply stringing these two functions together, we use a finalizer to ask the user for final confirmation before actually going through with the transaction:

```
getMoney :: Account → Int → IO ()
getMoney acc amount = atomicallyWithIO action finalizer
where
  action = do withdraw acc amount
          bal ← readTVar acc
          return bal
  finalizer bal = do putStrLn ("New balance: " ++ show bal)
                  putStrLn "Confirm (yes/no)?"
                  answer ← getLine
                  case answer of
                    "yes" → dispenseCash amount
                    _     → error "cancelled"
```

What this example shows is how we can use `atomicallyWithIO` to play out the revocable effects of a transaction (e.g. withdrawing money from an account) and use the results in a side-effecting manner (e.g. showing the user the new balance of her account and asking for confirmation) before committing to those effects.

Note that the balance we show to the user will be the final balance. There is no way another transaction could change it before we finish the I/O action. Once the I/O phase begins, the transaction is locked. At this point, only the finalizer itself can cause a rollback. And only when the finalizer finishes by returning is the transaction made visible to the rest of the world.

There are obvious performance implications with this scheme: if the I/O action takes a long time, other transactions will have to wait before they can make progress. This bottleneck is unavoidable, the transition from STM to I/O naturally creates a serialization point. It is up to the programmer to not waste undue time or otherwise ensure that the finalizer finishes in a timely manner, e.g. by using the standard `timeout` function.

2.2.3 Example 3: Nesting

Since the finalizer can be an arbitrary I/O action, the question arises: can we run `atomicallyWithIO` *inside* `atomicallyWithIO`?

Running `atomically` inside `atomically` is evidently not possible—the type system forbids it. Even something like

```
atomically (unsafeIOToSTM (atomically m))
```

will not work: GHC detects such nefariousness and throws a runtime exception. The runtime system does not support running a transaction inside another transaction, because there is no single intuitive way to resolve problems such as conflicting transactions or nested rollback. Luckily, it does not appear that this kind of nesting is necessary in practice or even useful.

But running `atomically` within a finalizer is a different story. First of all, this is eminently useful. Consider the following example, in which a counter is `atomically` increased only after confirmation has been received from two different servers. For efficiency, both servers are queried in parallel:

```
atomicallyWithIO (do n ← readTVar counter
                  writeTVar x (n + 1)
                  return n
                )
  (λn → do (r1, r2) ← concurrently (post n server1)
                                             (post n server2)
      if r1 ∧ r2
      then return ()
      else error "rejected"
  )
```

The function `concurrently :: IO a → IO b → IO (a, b)` is exported by the `async` library (Marlow 2013a), which provides “a set of operations for running IO operations asynchronously and waiting for their results.” The library is implemented using STM. When executing `concurrently`, there are actually multiple atomic transactions happening behind the scenes.

Many complex I/O actions might use STM internally.¹ There is no reason why these should not also be usable within the finalizer of `atomicallyWithIO`. The type system allows it, and the semantics are clear as well: finalizers are already regarded as irrevocable. A transaction run within the finalizer of another transaction is completely independent from that other transaction. There is only one restriction: *the inner transaction cannot modify any of the transactional variables used by the outer transaction*; this would inevitably lead to a deadlock, as the inner transaction would have to wait for the outer

¹ Indeed, as of GHC 7.8, the base library itself makes use of STM in `GHC.Event.Unique`. This module is in turn used by I/O primitives like `threadDelay`. Thus STM usage is actually spread throughout large parts of the I/O subsystem.

transaction to commit and release its locks on the shared variables, while the outer transaction can only commit once the inner transaction has committed.

But are these kinds of deadlock situations not exactly what STM is supposed to protect us from? True, but I feel that introducing the possibility of deadlocks—limited to this one specific scenario, in which the nested transactions operate on the same set of variables, which I believe is rather artificial—is justified by the great usefulness of allowing independent STM transactions to run during an STM finalizer. Also, the runtime system can always detect these kinds of deadlocks instead of simply looping forever, and it will throw an exception that gracefully aborts the transactions involved and that could be used to pinpoint the source of the bug.

In principle, circular dependencies between shared variables could even be detected statically, at compile time. Trying to enforce the non-circularity via the type system would probably involve advanced type system features implemented by GHC extensions, if not require dependent types outright. In any case, it would necessitate major backwards-incompatible changes to the existing STM API, which is something I wanted to avoid. Alternatively, one could imagine a separate pre-processing step in the compiler, but this seems rather inelegant and more trouble than it is worth.

We should note an important point here: a finalizer has the same global view of the world as any other I/O action running at the same time. In particular, the new value of a TVar updated during the STM part of the transaction will not be visible in the finalizer. The variable can be read safely—a deadlock only happens if an inner transaction tries to *modify* a shared variable—but the value returned will be its old value, which is the actual, globally known value of the TVar, at this moment in time.

2.3 Related Work

Extending Haskell’s STM to allow safe combination of atomic blocks with I/O actions has been proposed from the very beginning. In a 2006 post on the `haskell-cafe` mailing list, Simon Peyton Jones suggested an operator

$$\text{onCommit} :: \text{IO } a \rightarrow \text{STM } ()$$

that “would queue up an IO action to be performed when the transaction commits.”² The `stm-io-hooks` package (Robinson and Kuklewicz 2012) implements this operator in a custom STM monad that is meant as drop-in replacement for the system one. `onCommit` has the same semantics as `atomicallyWithIO`: “The commit IO action will be executed iff the transaction commits.” The difference is that `onCommit` is truly composable. Any STM function can use `onCommit` to add an I/O action to the list of actions to be executed when the atomic block commits. The caller of `atomically` does not have to be aware of it.

While greater composability seems to be more in the spirit of Haskell’s STM, I think that in this case it is actually dangerous. I/O is fundamentally not composable in the same way that STM is. Consider the following scenario:

²<http://www.haskell.org/pipermail/haskell-cafe/2006-November/019771.html>


```

foo :: STM ()
foo = do writeTVar a 1
      onCommit (serialize "a" 1)

bar :: STM ()
bar = onCommit (error "rollback")

baz :: IO ()
baz = atomically (foo >> bar)

```

Calling `baz` will lead to an inconsistent state: the `onCommit` action of `bar` will abort the whole transaction and so the STM effects of `foo` are rolled back, yet *the I/O effects are not*. Even worse, there is no way from looking only at `baz` to discern that any of this is going on. The composability of `onCommit` hides effects that in my opinion should be made explicit. This may not be an issue in all cases, but for serialization it clearly is. Using `atomicallyWithIO` is safer in this regard, at the cost of reduced composability.

2.4 Semantics of STM

I now formalize my design by giving an operational semantics. In this section, I present an overview of the original semantics of STM Haskell as described by Harris, Marlow, et al. (2005), including the revised exception semantics from their post-publication appendix. Using this as a foundation, I will then add the changes necessary to implement finalizers in Section 2.5. There I will also include the invariant semantics described by Harris and Peyton Jones (2006).

Throughout my presentation of these semantics, I aimed to keep as close to the original papers as possible to ease cross-referencing. However, a few minor modifications have been made to better reflect the current state of GHC, such as changing the names of some of the primitives.

Figures 2.1 to 2.4 give the complete syntax and semantics of the original version of STM Haskell. The key idea is to separate *I/O transitions* (“ \rightarrow ”, Figure 2.3) from *STM transitions* (“ \Rightarrow ”, Figure 2.4). Execution proceeds by non-deterministically choosing a thread and performing a single I/O transition. Thus, the execution of different threads may be interleaved and we have concurrency on the I/O level. However, STM transitions cannot interleave. They are only ever performed as premises of the `atomically` operator and the resulting state change is thus regarded as a single atomic step on the I/O level.

Simply requiring that `atomically` must reduce (on the STM level) to either `return` or `throw`, and not `retry`, obviates the need for modeling transaction logs or rollback. These matters are left for the implementation.

Syntax The syntax of values and terms (Figure 2.1) is entirely conventional, except that some monadic combinators are treated as values. The `do`-notation used in the preceding sections is of course just the usual syntactic sugar for the monadic operators `>>=` and `return`:

$$\begin{aligned} \text{do } \{ x \leftarrow e; Q \} &\equiv e \gg= (\backslash x \rightarrow \text{do } \{ Q \}) \\ \text{do } \{ e; Q \} &\equiv e \gg= (\backslash_ \rightarrow \text{do } \{ Q \}) \\ \text{do } \{ e \} &\equiv e \end{aligned}$$

A *program state* $P; \Theta$ consists of a *thread soup* P and a *heap* Θ . The thread soup P is a multi-set of threads with each thread consisting of a single term M_t , where t is the thread ID. The heap Θ maps references to terms.

Additionally, we sometimes want to explicitly track a set of *allocation effects* Δ . This is a redundant subset of the heap. We mainly need it during exception handling, where we want to roll back heap effects but have to retain allocation effects.

When describing a transition between program states, we use an *evaluation context* to identify the active site of the transition. The program evaluation context \mathbb{P} arbitrarily chooses a thread from the soup. This corresponds to the scheduler of a real implementation. STM terms are then evaluated using the context \mathbb{S} .

Administrative transitions A few fundamental transitions are used on both the I/O and the STM level (Figure 2.2). The rule (BIND), which implements sequential composition in the monad, as well as (THROW) and (RETRY) are quite ordinary. (EVAL) allows evaluation of terms. It uses a function \mathcal{V} , which is entirely standard and whose definition is omitted.

The I/O level rule (ADMIN) lifts administrative rules into the I/O world. (AADMIN) does the same for STM.

	x, y	\in	<i>Variable</i>
	r, t	\in	<i>Name</i>
	c	\in	<i>Char</i>
Value	V	$::=$	$r \mid c \mid \backslash x \rightarrow M$ $\mid \text{return } M \mid M \gg= N$ $\mid \text{putChar } c \mid \text{getChar}$ $\mid \text{throw } M \mid \text{catch } M N$ $\mid \text{retry} \mid M \text{ `orElse` } N$ $\mid \text{newTVar}$ $\mid \text{readTVar } r \mid \text{writeTVar } r M$ $\mid \text{forkIO } M$ $\mid \text{atomically } M$
Term	M, N	$::=$	$x \mid V \mid M N \mid \dots$
Thread soup	P, Q	$::=$	$M_t \mid (P \mid Q)$
Heap	Θ	$::=$	$r \hookrightarrow M$
Allocations	Δ	$::=$	$r \hookrightarrow M$
Evaluation contexts	\mathbb{S}	$::=$	$[\cdot] \mid \mathbb{S} \gg= M$
	\mathbb{P}	$::=$	$\mathbb{S}_t \mid (\mathbb{P} \mid P) \mid (P \mid \mathbb{P})$
Action	a	$::=$	$!c \mid ?c \mid \epsilon$

Figure 2.1: The syntax of STM Haskell (Harris, Marlow, et al. 2005)

Administrative transitions		$M \rightarrow N$
M	\rightarrow	V if $\mathcal{V}[\![M]\!] = V$ and $M \not\equiv V$ (EVAL)
$\text{return } N \gg= M$	\rightarrow	$M N$ (BIND)
$\text{throw } N \gg= M$	\rightarrow	$\text{throw } N$ (THROW)
$\text{retry} \gg= M$	\rightarrow	retry (RETRY)

Figure 2.2: Administrative transitions of STM Haskell (Harris, Marlow, et al. 2005)

I/O transitions A top level I/O transition (Figure 2.3) is of the form

$$P; \Theta \xrightarrow{a} Q; \Theta'.$$

It takes a program state $P; \Theta$ to a new state $Q; \Theta'$ while performing the input/output action a . The actions $!c$ and $?c$ denote writing a character to the standard output and reading a character from the standard input, respectively. The silent action ϵ does nothing and is usually omitted. These are the only actions in our model, a real system would of course have many more.

The first two rules, (PUTC) and (GETC), are applicable when the active term is `putChar` or `getChar`, respectively. The transition carries out the appropriate action and replaces the term with a `return` containing the result. The rule (FORK) creates a new thread, with a freshly chosen thread ID t , and adds it to the thread soup. (CATCH1) and (CATCH2) handle I/O level exceptions in the standard way.

More interesting are the rules (ARET) and (ATHROW). They define the semantics of atomic blocks. We can see that the only way to perform zero or more STM transitions (“ \Rightarrow ”) is by performing a single I/O transition (“ \rightarrow ”), brought about by the `atomically` keyword. Interleaving of STM transitions is therefore not possible. Also, these are the only I/O level rules that affect the heap, which means that the heap can only be mutated inside an atomic block.

The rule (ARET) concerns STM terms that ultimately evaluate to `return` statements. This means the STM transaction was successfully completed and the updated heap Θ' becomes visible on the I/O level. If, on the other hand, the STM term evaluates to `throw N`, i.e. if the STM transaction throws an exception N , the exception is propagated to the I/O level and the new heap discarded. Only the allocation effects Δ' are kept and made globally visible. This is necessary because the exception value N may contain references to variables that were allocated inside the transaction and we do not want to leave dangling pointers around.

I/O transitions			$P; \Theta \xrightarrow{a} Q; \Theta'$
$\mathbb{P}[\text{putChar } c]; \Theta$	$\xrightarrow{!c}$	$\mathbb{P}[\text{return } ()]; \Theta$	(PUTC)
$\mathbb{P}[\text{getChar}]; \Theta$	$\xrightarrow{?c}$	$\mathbb{P}[\text{return } c]; \Theta$	(GETC)
$\mathbb{P}[\text{forkIO } M]; \Theta$	\rightarrow	$(\mathbb{P}[\text{return } t] \mid M_t); \Theta \quad t \notin \mathbb{P}, \Theta, M$	(FORK)
$\mathbb{P}[\text{catch (return } M) N]; \Theta$	\rightarrow	$\mathbb{P}[\text{return } M]; \Theta$	(CATCH1)
$\mathbb{P}[\text{catch (throw } P) N]; \Theta$	\rightarrow	$\mathbb{P}[N P]; \Theta$	(CATCH2)
$\frac{M \rightarrow N}{\mathbb{P}[M]; \Theta \rightarrow \mathbb{P}[N]; \Theta} \quad (\text{ADMIN})$			
$\frac{M; \Theta, \{\} \xRightarrow{*} \text{return } N; \Theta', \Delta'}{\mathbb{P}[\text{atomically } M]; \Theta \rightarrow \mathbb{P}[\text{return } N]; \Theta'} \quad (\text{ARET})$			
$\frac{M; \Theta, \{\} \xRightarrow{*} \text{throw } N; \Theta', \Delta'}{\mathbb{P}[\text{atomically } M]; \Theta \rightarrow \mathbb{P}[\text{throw } N]; \Theta \cup \Delta'} \quad (\text{ATHROW})$			

Figure 2.3: I/O transitions of STM Haskell (Harris, Marlow, et al. 2005)

STM transitions An *STM transition* (Figure 2.4) is of the form

$$M; \Theta, \Delta \Rightarrow N; \Theta', \Delta'.$$

Θ is again the heap, while Δ redundantly records the allocation effects. Note that none of these effects are visible globally until we make a transition on the I/O level that exposes them. All STM transitions stay completely within the STM layer, except for rule (AADMIN), which just lifts pure administrative transitions.

Most of the STM rules are quite standard. (READ), (WRITE) and (NEW) describe the reading, writing and allocating of mutable variables. Note how (NEW) makes allocation effects explicit by modifying Δ .

The rules (OR1-3) handle the kind of nested transactions made possible by the `orElse` combinator.³ When encountering a term of the form $M_1 \text{ orElse } M_2$ we always begin by inspecting the left branch M_1 . If this evaluates to a `return` (OR1) or a `throw` (OR2), the result is propagated and all memory effects are retained. Keep in mind that we are still on the STM level. If the whole STM transaction finally evaluates to `throw`, then its memory effects are of course globally discarded by the I/O level rule (ATHROW). If we encounter a `retry` when evaluating M_1 , rule (OR3) applies, discarding all memory effects and continuing with the evaluation of the right branch M_2 .

Rules (XSTM1-3) describe exception handling within STM by use of the `catch` combinator. When evaluating the body M of a term `catch M N`, we start with the current heap Θ and an empty set of allocation effects. On successful evaluation (XSTM1), all effects are preserved and the catch handler N is ignored. If M throws an exception (XSTM2), that exception is given to the handler N and only the allocation effects of M are preserved. If M evaluates to `retry` (XSTM3), all effects are discarded.

This last set of rules was introduced in the post-publication appendix of Harris, Marlow, et al. (2005) to alleviate a subtle issue concerning leaking of effects. For a more elaborate discussion of this issue and of the whole system, refer to that paper.

³Not to be confused with the nesting of `atomically` calls described in Section 2.2.3.

STM transitions		$M; \Theta, \Delta \Rightarrow N; \Theta', \Delta'$
$\mathbb{S}[\text{readTVar } r]; \Theta, \Delta$	\Rightarrow	$\mathbb{S}[\text{return } \Theta(r)]; \Theta, \Delta$ if $r \in \text{dom}(\Theta)$ (READ)
$\mathbb{S}[\text{writeTVar } r M]; \Theta, \Delta$	\Rightarrow	$\mathbb{S}[\text{return } ()]; \Theta[r \mapsto M], \Delta$ if $r \in \text{dom}(\Theta)$ (WRITE)
$\mathbb{S}[\text{newTVar } M]; \Theta, \Delta$	\Rightarrow	$\mathbb{S}[\text{return } r]; \Theta[r \mapsto M], \Delta[r \mapsto M]$ $r \notin \text{dom}(\Theta)$ (NEW)
		$\frac{M \rightarrow N}{\mathbb{S}[M]; \Theta, \Delta \Rightarrow \mathbb{S}[N]; \Theta, \Delta} \text{ (AADMIN)}$
		$\frac{M_1; \Theta, \Delta \xRightarrow{*} \text{return } N; \Theta', \Delta'}{\mathbb{S}[M_1 \text{ `orElse` } M_2]; \Theta, \Delta \Rightarrow \mathbb{S}[\text{return } N]; \Theta', \Delta'} \text{ (OR1)}$
		$\frac{M_1; \Theta, \Delta \xRightarrow{*} \text{throw } N; \Theta', \Delta'}{\mathbb{S}[M_1 \text{ `orElse` } M_2]; \Theta, \Delta \Rightarrow \mathbb{S}[\text{throw } N]; \Theta', \Delta'} \text{ (OR2)}$
		$\frac{M_1; \Theta, \Delta \xRightarrow{*} \text{retry}; \Theta', \Delta'}{\mathbb{S}[M_1 \text{ `orElse` } M_2]; \Theta, \Delta \Rightarrow \mathbb{S}[M_2]; \Theta, \Delta} \text{ (OR3)}$
		$\frac{M; \Theta, \{\} \xRightarrow{*} \text{return } P; \Theta', \Delta'}{\mathbb{S}[\text{catch } M N]; \Theta, \Delta \Rightarrow \mathbb{S}[\text{return } P]; \Theta', (\Delta \cup \Delta')} \text{ (XSTM1)}$
		$\frac{M; \Theta, \{\} \xRightarrow{*} \text{throw } P; \Theta', \Delta'}{\mathbb{S}[\text{catch } M N]; \Theta, \Delta \Rightarrow \mathbb{S}[N P]; (\Theta \cup \Delta'), (\Delta \cup \Delta')} \text{ (XSTM2)}$
		$\frac{M; \Theta, \{\} \xRightarrow{*} \text{retry}; \Theta', \Delta'}{\mathbb{S}[\text{catch } M N]; \Theta, \Delta \Rightarrow \mathbb{S}[\text{retry}]; \Theta, \Delta} \text{ (XSTM3)}$

Figure 2.4: STM transitions of STM Haskell (Harris, Marlow, et al. 2005)

2.5 Semantics of finalizers

Now that we are familiar with the original STM semantics, we can extend them to support finalizers. First, let us realize that `atomically` is trivially expressed in terms of `atomicallyWithIO`:

$$\text{atomically } m \quad \equiv \quad \text{atomicallyWithIO } m \text{ return}$$

We can therefore use `atomicallyWithIO` to replace the old `atomically` completely. This keeps things simple and saves us from an annoying duplication of rules. Figure 2.5 gives the revised I/O transitions. There is no need to change anything on the STM level.

Rule (ATHROW) has been replaced by (ATHROW1), which is identical except that it now uses `atomicallyWithIO`. Like before, if the STM term evaluates to an exception, the changed heap Θ' is discarded and only the allocation effects Δ' are kept. The finalizer F given to `atomicallyWithIO` is simply ignored in this case.

The heart of the change is to be found in rule (ARET), which now has an additional premise: after the atomic block returns, the finalizer given to `atomicallyWithIO` must also evaluate to `return`. It does so by making zero or more *I/O-level* transitions. Note that these transitions begin with the original unchanged heap Θ . The finalizer executes in a global context, where the results of the STM transaction are not yet visible. Thus the finalizer cannot know about the modified heap Θ' . (It must know about the allocation effects Δ' , however: N could contain a reference to a newly allocated variable and we do not want any dangling pointers.) The final heap resulting from the (ARET) transition is the union of the STM result heap Θ' with the result heap of the I/O transition, $\hat{\Theta}$.

The new rule (ATHROW2) handles exceptions that occur during the coupled I/O transition. Like with (ATHROW1), the STM result heap Θ' is discarded. However, the I/O result heap $\hat{\Theta}$ *must be kept*, because it could have already been seen by other I/O level transitions. This is what we mean when we speak of irrevocable I/O.

I/O transitions $P; \Theta \xrightarrow{a} Q; \Theta'$		
$\mathbb{P}[\text{putChar } c]; \Theta$	$\xrightarrow{!c} \mathbb{P}[\text{return } ()]; \Theta$	(PUTC)
$\mathbb{P}[\text{getChar}]; \Theta$	$\xrightarrow{?c} \mathbb{P}[\text{return } c]; \Theta$	(GETC)
$\mathbb{P}[\text{forkIO } M]; \Theta$	$\rightarrow (\mathbb{P}[\text{return } t] \mid M_t); \Theta \quad t \notin \mathbb{P}, \Theta, M$	(FORK)
$\mathbb{P}[\text{catch (return } M) N]; \Theta$	$\rightarrow \mathbb{P}[\text{return } M]; \Theta$	(CATCH1)
$\mathbb{P}[\text{catch (throw } P) N]; \Theta$	$\rightarrow \mathbb{P}[N P]; \Theta$	(CATCH2)
$\frac{M \rightarrow N}{\mathbb{P}[M]; \Theta \rightarrow \mathbb{P}[N]; \Theta} \quad (\text{ADMIN})$		
$\frac{\begin{array}{l} M; \Theta, \{\} \xRightarrow{*} \text{return } N; \Theta', \Delta' \\ F N; (\Theta \cup \Delta') \xrightarrow{*} \text{return } P; \hat{\Theta} \end{array}}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta \rightarrow \mathbb{P}[\text{return } P]; \Theta' \cup \hat{\Theta}} \quad (\text{ARET})$		
$\frac{M; \Theta, \{\} \xRightarrow{*} \text{throw } N; \Theta', \Delta'}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta \rightarrow \mathbb{P}[\text{throw } N]; \Theta \cup \Delta'} \quad (\text{ATHROW1})$		
$\frac{\begin{array}{l} M; \Theta, \{\} \xRightarrow{*} \text{return } N; \Theta', \Delta' \\ F N; (\Theta \cup \Delta') \xrightarrow{*} \text{throw } P; \hat{\Theta} \end{array}}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta \rightarrow \mathbb{P}[\text{return } P]; \hat{\Theta}} \quad (\text{ATHROW2})$		

Figure 2.5: I/O transitions with finalizers

2.5.1 Nesting

The attentive reader might have noticed that in the extended semantics given above the following term is perfectly valid:

$$\text{atomicallyWithIO } M_1 (\text{atomicallyWithIO } M_2 F)$$

As I have discussed in Section 2.2, this kind of nesting is actually pretty useful in practice. However, it does open up the possibility of deadlocks if the term M_2 writes to any of the memory locations referenced by M_1 . Such behavior should be explicitly prohibited in the semantics, which need to be extended a bit further to do so (Figures 2.6 and 2.7).

Firstly, both I/O and STM transitions will carry around a new piece of state Σ , which is a set of labeled memory references:

$$\Sigma ::= r \times (R \mid W)$$

It is a record of all memory operations done by a transaction. The labels R and W denote whether a particular memory location was read from or written to, respectively. As a shorthand notation, we define $\Sigma_W = \{r \mid (r, W) \in \Sigma\}$, i.e. the set of references in Σ that have been used in a write operation. Σ_R and Σ_{RW} are defined analogously. We also define $\Sigma|_\Delta = \{(r, w) \in \Sigma \mid r \in \text{dom}(\Delta)\}$, which restricts Σ to those labelled references that correspond to the allocations recorded in Δ .

Secondly, the STM transition rules (READ), (WRITE) and (NEW) modify Σ by recording their respective operations. All other rules treat Σ in very much the same way as the heap.

Finally, the purpose of Σ becomes clear on the I/O-level, where rules (ARET) and (ATHROW2) include a new condition above the line: the set of memory locations read from or written to during the STM transition (Σ'_{RW}) and the set of memory locations written to during the finalizing I/O transition ($\hat{\Sigma}_W$) must be entirely distinct. There must not be a single shared reference between the two phases—except for references that the finalizer has only read. This is exactly the restriction necessary to avoid deadlocks.

STM transitions		$M; \Theta, \Delta, \Sigma \Rightarrow N; \Theta', \Delta', \Sigma'$
$\mathbb{S}[\text{readTVar } r]; \Theta, \Delta, \Sigma$	\Rightarrow	$\mathbb{S}[\text{return } \Theta(r)]; \Theta, \Delta, (\Sigma \cup \{(r, \mathbf{R})\})$ if $r \in \text{dom}(\Theta)$ (READ)
$\mathbb{S}[\text{writeTVar } r \ M]; \Theta, \Delta, \Sigma$	\Rightarrow	$\mathbb{S}[\text{return } ()]; \Theta[r \mapsto M], \Delta, (\Sigma \cup \{(r, \mathbf{W})\})$ if $r \in \text{dom}(\Theta)$ (WRITE)
$\mathbb{S}[\text{newTVar } M]; \Theta, \Delta, \Sigma$	\Rightarrow	$\mathbb{S}[\text{return } r]; \Theta[r \mapsto M], \Delta[r \mapsto M], (\Sigma \cup \{(r, \mathbf{W})\})$ $r \notin \text{dom}(\Theta)$ (NEW)
$\frac{M \rightarrow N}{\mathbb{S}[M]; \Theta, \Delta, \Sigma \Rightarrow \mathbb{S}[N]; \Theta, \Delta, \Sigma} \text{ (AADMIN)}$		
$\frac{M_1; \Theta, \Delta, \Sigma \xRightarrow{*} \text{return } N; \Theta', \Delta', \Sigma'}{\mathbb{S}[M_1 \text{ `orElse` } M_2]; \Theta, \Delta, \Sigma \Rightarrow \mathbb{S}[\text{return } N]; \Theta', \Delta', \Sigma'} \text{ (OR1)}$		
$\frac{M_1; \Theta, \Delta, \Sigma \xRightarrow{*} \text{throw } N; \Theta', \Delta', \Sigma'}{\mathbb{S}[M_1 \text{ orElse } M_2]; \Theta, \Delta, \Sigma \Rightarrow \mathbb{S}[\text{throw } N]; \Theta', \Delta', \Sigma'} \text{ (OR2)}$		
$\frac{M_1; \Theta, \Delta, \Sigma \xRightarrow{*} \text{retry}; \Theta', \Delta', \Sigma'}{\mathbb{S}[M_1 \text{ orElse } M_2]; \Theta, \Delta, \Sigma \Rightarrow \mathbb{S}[M_2]; \Theta, \Delta, \Sigma} \text{ (OR3)}$		
$\frac{M; \Theta, \{\}, \Sigma \xRightarrow{*} \text{return } P; \Theta', \Delta', \Sigma'}{\mathbb{S}[\text{catch } M \ N]; \Theta, \Delta, \Sigma \Rightarrow \mathbb{S}[\text{return } P]; \Theta', (\Delta \cup \Delta'), \Sigma'} \text{ (XSTM1)}$		
$\frac{M; \Theta, \{\}, \Sigma \xRightarrow{*} \text{throw } P; \Theta', \Delta', \Sigma'}{\mathbb{S}[\text{catch } M \ N]; \Theta, \Delta, \Sigma \Rightarrow \mathbb{S}[N \ P]; (\Theta \cup \Delta'), (\Delta \cup \Delta'), (\Sigma \cup \Sigma' _{\Delta'})} \text{ (XSTM2)}$		
$\frac{M; \Theta, \{\}, \Sigma \xRightarrow{*} \text{retry}; \Theta', \Delta', \Sigma'}{\mathbb{S}[\text{catch } M \ N]; \Theta, \Delta, \Sigma \Rightarrow \mathbb{S}[\text{retry}]; \Theta, \Delta, \Sigma} \text{ (XSTM3)}$		

Figure 2.6: STM transitions with **finalizers** and **nesting checks**

I/O transitions			
	$P; \Theta, \Sigma \xrightarrow{a} Q; \Theta', \Sigma'$		
$\mathbb{P}[\text{putChar } c]; \Theta, \Sigma$	$\xrightarrow{!c} \mathbb{P}[\text{return } ()]; \Theta, \Sigma$	(PUTC)	
$\mathbb{P}[\text{getChar}]; \Theta, \Sigma$	$\xrightarrow{?c} \mathbb{P}[\text{return } c]; \Theta, \Sigma$	(GETC)	
$\mathbb{P}[\text{forkIO } M]; \Theta, \Sigma$	$\rightarrow (\mathbb{P}[\text{return } t] \mid M_t); \Theta, \Sigma \quad t \notin \mathbb{P}, \Theta, \Sigma, M$	(FORK)	
$\mathbb{P}[\text{catch (return } M) N]; \Theta, \Sigma$	$\rightarrow \mathbb{P}[\text{return } M]; \Theta, \Sigma$	(CATCH1)	
$\mathbb{P}[\text{catch (throw } P) N]; \Theta, \Sigma$	$\rightarrow \mathbb{P}[N P]; \Theta, \Sigma$	(CATCH2)	
$\frac{M \rightarrow N}{\mathbb{P}[M]; \Theta, \Sigma \rightarrow \mathbb{P}[N]; \Theta, \Sigma} \quad (\text{ADMIN})$			
$\frac{\begin{array}{l} M; \Theta, \{\}, \{\} \xRightarrow{*} \text{return } N; \Theta', \Delta', \Sigma' \\ F N; (\Theta \cup \Delta'), \{\} \xrightarrow{*} \text{return } P; \hat{\Theta}, \hat{\Sigma} \quad (\Sigma'_{\text{RW}} \cap \hat{\Sigma}_{\text{W}}) = \emptyset \end{array}}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta, \Sigma \rightarrow \mathbb{P}[\text{return } P]; (\Theta' \cup \hat{\Theta}), (\Sigma \cup \Sigma' \cup \hat{\Sigma})} \quad (\text{ARET})$			
$\frac{M; \Theta, \{\}, \{\} \xRightarrow{*} \text{throw } N; \Theta', \Delta', \Sigma'}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta, \Sigma \rightarrow \mathbb{P}[\text{throw } N]; (\Theta \cup \Delta'), (\Sigma \cup \Sigma' _{\Delta'})} \quad (\text{ATHROW1})$			
$\frac{\begin{array}{l} M; \Theta, \{\}, \{\} \xRightarrow{*} \text{return } N; \Theta', \Delta', \Sigma' \\ F N; (\Theta \cup \Delta'), \{\} \xrightarrow{*} \text{throw } P; \hat{\Theta}, \hat{\Sigma} \quad (\Sigma'_{\text{RW}} \cap \hat{\Sigma}_{\text{W}}) = \emptyset \end{array}}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta, \Sigma \rightarrow \mathbb{P}[\text{throw } P]; \hat{\Theta}, (\Sigma \cup \Sigma' _{\Delta'} \cup \hat{\Sigma})} \quad (\text{ATHROW2})$			

Figure 2.7: I/O transitions with **finalizers** and **nesting checks**

2.5.2 Adding invariants

GHC’s STM implementation includes the `alwaysSucceeds` primitive, used to introduce dynamically-checked data invariants. It was first described in Harris and Peyton Jones (2006) (where it was called `check`). I have ignored invariants thus far since they increase the overall complexity of the semantics and are mostly orthogonal to finalizers. It is now time to introduce them. Figures 2.8 to 2.10 constitute the complete and final syntax and semantics of STM Haskell with finalizers.

Firstly, we add the set of invariants Ω as another piece of state associated with all I/O and STM transitions. It is usually treated like Σ or the heap.

Secondly, the new STM transition rules (CHECK1) and (CHECK2) describe the `alwaysSucceeds` primitive. If the invariant holds at the point it is proposed, (CHECK1) adds it to Ω and discards its evaluation effects. (CHECK2) applies when the invariant throws an exception. As usual, the exception is propagated, the invariant discarded and only allocation effects are retained.

Finally, the biggest changes happen on the I/O level: (ARET) becomes (ARET1); along with (ATHROW2) it has a new premise that evaluates the invariants in place at the end of the STM transition. When evaluating each invariant, its return value and all its heap effects are discarded. Only if all invariants evaluate to `return` terms can we run the finalizer of `atomicallyWithIO`. Note that it is required that the write set of the I/O transition ($\hat{\Sigma}_w$) is distinct from all of the memory locations referenced by the invariants ($\bigcup \Sigma'_i$).

If any of the invariants throws an exception, the new rule (ARET2) applies. Similar to (ATHROW1), the exception is propagated to the I/O level and only allocation effects are preserved (including the allocation effects Δ'_i of the broken invariant M_i). The I/O action is not executed in this case.

	x, y	\in	<i>Variable</i>
	r, t	\in	<i>Name</i>
	c	\in	<i>Char</i>
Value	V	$::=$	$r \mid c \mid \backslash x \rightarrow M$ \mid $\text{return } M \mid M \gg= N$ \mid $\text{putChar } c \mid \text{getChar}$ \mid $\text{throw } M \mid \text{catch } M N$ \mid $\text{retry} \mid M \text{ `orElse` } N$ \mid newTVar \mid $\text{readTVar } r \mid \text{writeTVar } r M$ \mid $\text{forkIO } M$ \mid $\text{alwaysSucceeds } M$ \mid $\text{atomicallyWithIO } M X$
Term	M, N	$::=$	$x \mid V \mid M N \mid \dots$
Thread soup	P, Q	$::=$	$M_t \mid (P \mid Q)$
Heap	Θ	$::=$	$r \hookrightarrow M$
Allocations	Δ	$::=$	$r \hookrightarrow M$
Touched references	Σ	$::=$	$r \times (\text{R} \mid \text{W})$
Invariants	Ω	$::=$	$\{M\}$
Evaluation contexts	\mathbb{S}	$::=$	$[\cdot] \mid \mathbb{S} \gg= M$
	\mathbb{P}	$::=$	$\mathbb{S}_t \mid (\mathbb{P} \mid P) \mid (P \mid \mathbb{P})$
Action	a	$::=$	$!c \mid ?c \mid \epsilon$

Figure 2.8: The syntax of STM Haskell with invariants, finalizers and nesting checks

I/O transitions			$P; \Theta, \Sigma, \Omega \xrightarrow{a} Q; \Theta', \Sigma', \Omega'$
$\mathbb{P}[\text{putChar } c]; \Theta, \Sigma, \Omega$	$\xrightarrow{!c}$	$\mathbb{P}[\text{return } ()]; \Theta, \Sigma, \Omega$	(PUTC)
$\mathbb{P}[\text{getChar}]; \Theta, \Sigma, \Omega$	$\xrightarrow{?c}$	$\mathbb{P}[\text{return } c]; \Theta, \Sigma, \Omega$	(GETC)
$\mathbb{P}[\text{forkIO } M]; \Theta, \Sigma, \Omega$	\rightarrow	$(\mathbb{P}[\text{return } t] \mid M_t); \Theta, \Sigma, \Omega \quad t \notin \mathbb{P}, \Theta, \Sigma, \Omega, M$	(FORK)
$\mathbb{P}[\text{catch (return } M) N]; \Theta, \Sigma, \Omega$	\rightarrow	$\mathbb{P}[\text{return } M]; \Theta, \Sigma, \Omega$	(CATCH1)
$\mathbb{P}[\text{catch (throw } P) N]; \Theta, \Sigma, \Omega$	\rightarrow	$\mathbb{P}[N P]; \Theta, \Sigma, \Omega$	(CATCH2)
$\frac{M \rightarrow N}{\mathbb{P}[M]; \Theta, \Sigma, \Omega \rightarrow \mathbb{P}[N]; \Theta, \Sigma, \Omega} \quad (\text{ADMIN})$			
$\frac{\begin{array}{l} M; \Theta, \{\}, \{\}, \Omega \xRightarrow{*} \text{return } N; \Theta', \Delta', \Sigma', \Omega' \\ \forall M_i \in \Omega' : (M_i; \Theta', \{\}, \{\}, \{\} \xRightarrow{*} \text{return } N_i; \Theta'_i, \Delta'_i, \Sigma'_i, \Omega'_i) \\ F N; (\Theta \cup \Delta'), \{\}, \Omega \xRightarrow{*} \text{return } P; \hat{\Theta}, \hat{\Sigma}, \hat{\Omega} \quad (\bigcup \Sigma'_i \cup \Sigma')_{\text{RW}} \cap \hat{\Sigma}_{\text{W}} = \emptyset \end{array}}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta, \Sigma, \Omega \rightarrow \mathbb{P}[\text{return } P]; (\Theta' \cup \hat{\Theta}), (\Sigma \cup \Sigma' \cup \hat{\Sigma}), (\Omega' \cup \hat{\Omega})} \quad (\text{ARET1})$			
$\frac{\begin{array}{l} M; \Theta, \{\}, \{\}, \Omega \xRightarrow{*} \text{return } N; \Theta', \Delta', \Sigma', \Omega' \\ \forall M_i \in \Omega' : (M_i; \Theta', \{\}, \{\}, \{\} \xRightarrow{*} \text{return } N_i; \Theta'_i, \Delta'_i, \Sigma'_i, \Omega'_i) \\ F N; (\Theta \cup \Delta'), \{\}, \Omega \xRightarrow{*} \text{throw } P; \hat{\Theta}, \hat{\Sigma}, \hat{\Omega} \quad (\bigcup \Sigma'_i \cup \Sigma')_{\text{RW}} \cap \hat{\Sigma}_{\text{W}} = \emptyset \end{array}}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta, \Sigma, \Omega \rightarrow \mathbb{P}[\text{throw } P]; \hat{\Theta}, (\Sigma \cup \Sigma' _{\Delta'} \cup \hat{\Sigma}), \hat{\Omega}} \quad (\text{ATHROW2})$			
$\frac{\begin{array}{l} M; \Theta, \{\}, \{\}, \Omega \xRightarrow{*} \text{return } N; \Theta', \Delta', \Sigma', \Omega' \\ \exists M_i \in \Omega' : (M_i; \Theta', \{\}, \{\}, \{\} \xRightarrow{*} \text{throw } N_i; \Theta'_i, \Delta'_i, \Sigma'_i, \Omega'_i) \end{array}}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta, \Sigma, \Omega \rightarrow \mathbb{P}[\text{throw } N_i]; (\Theta \cup \Delta' \cup \Delta'_i), (\Sigma \cup \Sigma' _{\Delta'} \cup \Sigma'_i _{\Delta'_i}), \Omega} \quad (\text{ARET2})$			
$\frac{M; \Theta, \{\}, \{\}, \Omega \xRightarrow{*} \text{throw } N; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{P}[\text{atomicallyWithIO } M F]; \Theta, \Sigma, \Omega \rightarrow \mathbb{P}[\text{throw } N]; (\Theta \cup \Delta'), (\Sigma \cup \Sigma' _{\Delta'}), \Omega} \quad (\text{ATHROW1})$			

Figure 2.9: I/O transitions with invariants, finalizers and nesting checks

STM transitions		$M; \Theta, \Delta, \Sigma, \Omega \Rightarrow N; \Theta', \Delta', \Sigma', \Omega'$	
$\mathbb{S}[\text{readTVar } r]; \Theta, \Delta, \Sigma, \Omega$	\Rightarrow	$\mathbb{S}[\text{return } \Theta(r)]; \Theta, \Delta, (\Sigma \cup \{(r, \mathbb{R})\}), \Omega$	if $r \in \text{dom}(\Theta)$ (READ)
$\mathbb{S}[\text{writeTVar } r \ M]; \Theta, \Delta, \Sigma, \Omega$	\Rightarrow	$\mathbb{S}[\text{return } ()]; \Theta[r \mapsto M], \Delta, (\Sigma \cup \{(r, w)\}), \Omega$	if $r \in \text{dom}(\Theta)$ (WRITE)
$\mathbb{S}[\text{newTVar } M]; \Theta, \Delta, \Sigma, \Omega$	\Rightarrow	$\mathbb{S}[\text{return } r]; \Theta[r \mapsto M], \Delta[r \mapsto M], (\Sigma \cup \{(r, w)\}), \Omega$	$r \notin \text{dom}(\Theta)$ (NEW)
$\frac{M \rightarrow N}{\mathbb{S}[M]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[N]; \Theta, \Delta, \Sigma, \Omega} \text{ (AADMIN)}$			
$\frac{M_1; \Theta, \Delta, \Sigma, \Omega \xRightarrow{*} \text{return } N; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{S}[M_1 \text{ `orElse` } M_2]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[\text{return } N]; \Theta', \Delta', \Sigma', \Omega'} \text{ (OR1)}$			
$\frac{M_1; \Theta, \Delta, \Sigma, \Omega \xRightarrow{*} \text{throw } N; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{S}[M_1 \text{ `orElse` } M_2]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[\text{throw } N]; \Theta', \Delta', \Sigma', \Omega'} \text{ (OR2)}$			
$\frac{M_1; \Theta, \Delta, \Sigma, \Omega \xRightarrow{*} \text{retry}; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{S}[M_1 \text{ `orElse` } M_2]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[M_2]; \Theta, \Delta, \Sigma, \Omega} \text{ (OR3)}$			
$\frac{M; \Theta, \{\}, \Sigma, \Omega \xRightarrow{*} \text{return } P; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{S}[\text{catch } M \ N]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[\text{return } P]; \Theta', (\Delta \cup \Delta'), \Sigma', \Omega'} \text{ (XSTM1)}$			
$\frac{M; \Theta, \{\}, \Sigma, \Omega \xRightarrow{*} \text{throw } P; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{S}[\text{catch } M \ N]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[N \ P]; (\Theta \cup \Delta'), (\Delta \cup \Delta'), (\Sigma \cup \Sigma' _{\Delta'}), \Omega} \text{ (XSTM2)}$			
$\frac{M; \Theta, \{\}, \Sigma, \Omega \xRightarrow{*} \text{retry}; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{S}[\text{catch } M \ N]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[\text{retry}]; \Theta, \Delta, \Sigma, \Omega} \text{ (XSTM3)}$			
$\frac{M; \Theta, \{\}, \Sigma, \Omega \xRightarrow{*} \text{return } N; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{S}[\text{alwaysSucceeds } M]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[\text{return } ()]; \Theta, \Delta, \Sigma, (\Omega \cup \{M\})} \text{ (CHECK1)}$			
$\frac{M; \Theta, \{\}, \Sigma, \Omega \xRightarrow{*} \text{throw } N; \Theta', \Delta', \Sigma', \Omega'}{\mathbb{S}[\text{alwaysSucceeds } M]; \Theta, \Delta, \Sigma, \Omega \Rightarrow \mathbb{S}[\text{throw } N]; (\Theta \cup \Delta'), (\Delta \cup \Delta'), (\Sigma \cup \Sigma' _{\Delta'}), \Omega} \text{ (CHECK2)}$			

Figure 2.10: STM transitions with invariants, finalizers and nesting checks

2.6 Implementation

The changes necessary to add finalizers to GHC’s STM implementation are surprisingly few. Before describing them I will once again first give a brief overview of the status quo. To avoid getting bogged down in minutiae, my description of the existing implementation will be a careful simplification, focusing only on the relevant parts of the system. The interested reader is referred to the STM Commentary (Yates 2013) for a more thorough description, and to the GHC source code itself for all the gory details.⁴ A fork of GHC containing the changes described in this section is available at <http://github.com/mcschroeder/ghc>.

2.6.1 Original STM interface

When evaluating the `atomically` function, the runtime system pushes an `ATOMICALLY_FRAME` onto the execution stack of the current thread. It then calls `stmStartTransaction` to initialize a new thread-local *transactional record* (TRec). During the transaction, all read and write operations are recorded in the TRec. Every TVar accessed by the transaction has an entry in the record. Each entry contains a reference to the value of the TVar at the start of the transaction (`expected_value`) and a reference to the new value the TVar will have when the transaction commits (`new_value`). If a TVar has only been read during a transaction, these two values will be identical. The TVars themselves are not modified until the transaction commits.

When execution returns to the `ATOMICALLY_FRAME`, the commit is initiated by calling `stmCommitTransaction`. First, the TRec is *validated* by checking if the `expected_value` of each entry is pointer-equal to the `current_value` field of the corresponding TVar, i.e. if the TVar has changed in between the transaction starting and committing. If everything is as expected, the TVar is locked by setting `current_value` to point to the committing TRec. If there is a mismatch, i.e. the TVar was modified by someone else, the TRec is discarded and the transaction restarted. When validation is successful for the whole TRec, the TVars are one by one updated by setting their `current_value` to the `new_value` of the corresponding TRec entry, which also unlocks them again.

Because all TVars are locked during validation and each one is only unlocked after it has been updated, the whole commit happens atomically with respect to other transactions. Any other transaction trying to commit at the same time will fail its validation. A transaction trying to read a locked TVar will briefly block, but since TVars are only ever locked for very short periods of time, this is not a problem.

2.6.2 Adding `atomicallyWithIO`

The finalizer given to `atomicallyWithIO` should be executed after it is guaranteed that the transaction can no longer abort due to conflicts with other transactions, but before the transaction’s memory effects are made visible. This is only possible at one mo-

⁴<http://git.haskell.org/ghc.git>

```

// Basic transaction execution
TRec *stmStartTransaction();
Closure *stmReadTVar(TRec *trec, TVar *tvar);
void stmWriteTVar(TRec *trec, TVar *tvar, Closure *new_value);

// Transaction commit operations
Bool stmPrepareToCommitTransaction(TRec *trec);
void stmCommitTransaction(TRec *trec);

// Blocking operations
Bool stmWait(TRec *trec);
Bool stmReWait(TRec *trec);

// Transactional variable
struct TVar {
    Closure    *current_value;
    TRec       *frozen_by;
}

// Transactional record entry
struct TRecEntry {
    TVar        *tvar;
    Closure     *expected_value;
    Closure     *new_value;
    TRecEntry   *next_entry;
}

// Transactional record
struct TRec {
    TRec        *enclosing_trec;
    TRecEntry   *next_entry;
}

```

Figure 2.11: STM runtime interface (simplified), extended to support [finalizers](#)

ment: during `stmCommitTransaction`, between the successful validation of the `TRec` and the updating of the `TVars`, while all the `TVars` are still locked.

To achieve this, the validation part of `stmCommitTransaction` has been split off into the new function `stmPrepareToCommitTransaction`. Committing is now a three-part sequence: first, `stmPrepareToCommitTransaction` is called to validate the transactions; if successful, a new `ATOMICALLY_FRAME` is pushed onto the stack and execution jumps to the finalizer code; when the finalizer is done and execution has returned to the `ATOMICALLY_FRAME`, `stmCommitTransaction` is called to finish the commit and update the `TVars`.

There is one problem: if the `TVars` are locked while the finalizer is running, then any transaction attempting to access any of those locked `TVars` will immediately block until the finalizer is finished. This is obviously bad for performance, especially considering that the blocking occurs even when both transactions are read-only. Giving up on read-parallelism is clearly unacceptable. But not locking the `TVars` while the finalizer is running is also not possible: the transaction could then unknowingly become invalidated by the actions of another transaction and still commit, resulting in an inconsistent state.

The solution: we *freeze* the `TVars` during finalization. This puts them into a kind of in-between state, where they can still be read by other transactions, but any transaction that attempts to commit an update to a frozen `TVar` will block until the `TVar` is unfrozen again. The new `TVar` field `frozen_by` is set during `stmPrepareToCommitTransaction`, after all `TVars` have been successfully validated, and points to the finalizing `TRec`; `frozen_by` is cleared again during `stmCommitTransaction`. Although the `TVars` need to be locked briefly when modifying their `frozen_by` fields (or indeed any other of their fields), they no longer need to be locked for the whole duration of the finalizer.

Unfortunately, it is necessary to lock *all* `TVars` that are involved in a transaction, *if* the transaction does indeed have a finalizer. Usually a read-phase is used, meaning `TVars` that are only read but not updated within the atomic block do not get locked at all, not even briefly. This is merely a performance optimization. It would be possible to do it in the presence of finalizers, but it complicates the implementation, since the `frozen_by` field is protected by the `TVar` lock. Note that this read-phase is also disabled in the original implementation whenever a transaction touches an invariant, for the very same reason.

What exactly happens when a transaction tries to commit an update but hits upon a frozen `TVar`? The transaction could simply restart, like it does whenever the `current_value` of a `TVar` changes behind its back. However, if the `TVar` is frozen, that means it is involved in a potentially long-running finalizer, so it could still be frozen when it is time to commit again. Since the transaction can never make progress as long as the `TVar` is frozen, the only logical solution is to wait until the state of the `TVar` has actually changed.⁵ This is similar to what happens when calling the blocking operator `retry`. In fact, we can reuse a lot of `retry`'s runtime machinery: if `stmPrepareToCommitTransaction`

⁵Additionally, there is an issue with GHC's optimizer being too eager when removing yield points. The transaction would just restart over and over again, denying the thread running the finalizer the chance to actually complete. This is a long-standing problem in GHC. See <http://ghc.haskell.org/trac/ghc/ticket/367>.

discovers that some of the transaction's `TVars` are frozen, it calls `stmWait` to put the `TRec` on the `TVars`' watch queues before putting the thread to sleep. Being on the `TVars`' watch queues means that whenever one of those `TVars` is updated in a commit, the waiting `TRec` is woken up. When this happens, `stmReWait` is called to validate the `TRec` again to see if it should continue waiting or if it makes sense now to restart the transaction or even to commit it right away.

2.6.3 Nesting

The original STM code was written under the assumption that nested transactions are impossible. To ensure safety, any attempts to bypass the type system are caught at run-time. For example, making calls of the following form will result in an error:

```
atomically (unsafeIOToSTM (atomically...))
```

However, the system does implicitly support two other kinds of nesting: (1) the `orElse` operator begins a new `TRec` for each of its branches; if a branch completes successfully (i.e. does not retry) the `TRec` is merged upwards into the enclosing record; (2) invariants introduced by `alwaysSucceeds` are checked at commit time, where a new nested transaction is created for each invariant and completely discarded once the check is over.

The form of nesting made possible by `atomicallyWithIO`, viz. beginning a new transaction within the finalizer of another transaction, is similar in execution to these two, except that we want to neither merge nor discard the nested `TRec`, but simply commit it as-is and then return to the old transaction. The only difference between a regular atomic block and an atomic block begun inside a finalizer is that in the latter case the thread has a dormant `TRec` it returns to once the transaction inside the finalizer has finished. Supporting this is pretty easy, as the existing infrastructure for nested `TRecs`, mainly the `enclosing_trec` field, can be reused. Supporting nested transactions basically amounts to loosening some assertions about what constitutes an outermost transaction in a chain of `TRecs`: previously, the outermost transaction was indicated by an empty `enclosing_trec` field, while now it can also be a transaction whose `enclosing_trec` is running a finalizer. Everything else works exactly as it did before.

What happens when the nested transactions share variables? The inner transaction just reading from a `TVar` is never a problem. It simply sees the `current_value` of the `TVar` in memory. Note that since the outer transaction has not yet committed, any updates it might have made to this `TVar` are still only in its `TRec`. They are not yet globally visible, including within the transaction's own finalizer. This matches the semantics. As does the fact that an inner transaction cannot write to a shared `TVar`. If it tries to, it will block during commit, since the `TVar` has been frozen by the outer transaction. Meanwhile, the outer transaction is waiting for the inner transaction to finish. The runtime system immediately detects this deadlock and throws an exception, gracefully aborting the whole transaction.

Transactional Tries

The fundamental data type of STM is the transactional variable. A `TVar` stores arbitrary data, to be accessed and modified in a thread-safe manner. For example, I might define a bank account as

```
type Euro = Int
type Account = TVar Euro
```

and then use a function like

```
transfer :: Account → Account → Euro → STM ()
```

to safely — in the transactional sense — move money between accounts.

But where do those accounts come from? If I am a bank, how do I represent the whole collection of accounts I manage, in a way that is transactionally safe? The obvious solution, and a common pattern, is to simply use an existing container type and put that type into a `TVar`:

```
type IBAN = String
type Bank = TVar (Map IBAN Account)
```

Since looking up an account from the `Map` involves a `readTVar` operation, the `Map` is entangled with the transaction, and I can be sure that when transferring money between accounts, both accounts actually exist in the bank at the time when the transaction commits.

The drawback of this pattern of simply wrapping a `Map` inside a `TVar` is that when adding or removing elements of the `Map`, one has to replace the `Map` inside the `TVar` wholesale. Thus all concurrently running transactions that have accessed the `Map` become invalid and will have to restart once they try to commit. Depending on the exact access patterns, this can be a serious cause of contention. For example, one benchmark

running on a 16-core machine, with 16 threads each trying to commit a slice out of 200 000 randomly generated transactions, resulted in over 1.3 million retries.¹ That is some serious overhead!

The underlying problem is that the whole `Map` is made transactional, when we only ever care about the subset of the `Map` that is relevant to the current transaction. If transaction *A* updates an element with key k_1 and transaction *B* deletes an element with key k_2 , then those two transactions only conflict if $k_1 = k_2$; if k_1 and k_2 are different, then there is no reason for either of the transactions to wait for the other one. But the `Map` does not know it is part of a transaction, and the `TVar` does not know nor care about the structure of its contents. And so the transactional net is cast too wide.

The solution is to not simply put a `Map`, or any other ready-made container type, into a `TVar`, but to design data structures specifically tailored to the needs of transactional concurrency. In this chapter, I present one such data structure, the *transactional trie*, based on the concurrent trie of Prokopec, Bagwell, and Odersky (2011). I will describe its design, show major parts of its implementation in Haskell and discuss the trade-offs that have to be made in the transactional setting. I will then evaluate it against similar STM-specialized data structures.

The finalizers of Chapter 2 were a macro-level approach to add side-effects to STM, connecting whole transactions with potentially large I/O actions, tailored to the use case of serializing data. The transactional trie incorporates side-effects on the micro-level. It exposes a transactionally safe interface, while internally circumventing some transactional safety measures in a controlled and manually verified way.

3.1 Background

The transactional trie is based on the concurrent trie of Prokopec, Bagwell, and Odersky (2011), which is a non-blocking concurrent version of the hash array mapped trie first described by Bagwell (2001).

A hash array mapped trie is a tree whose leaves store key-value bindings and whose nodes are implemented as arrays. Each array has 2^k elements. To look up a key, you take the initial k bits of the key's hash as an index into the root array. If the element at that index is another array node, you continue by using the next k bits of the hash as an index into that second array. If that element is another array, you again use the next k bits of the hash, and so on. Generally speaking, to index into an array node at level l , you use the k bits of the hash beginning at position $k * l$. This procedure is repeated until either a leaf node is found or one of the array nodes does not have an entry at the particular index, in which case the key is not yet present in the trie. The expected depth of the trie is $O(\log_{2^k}(n))$, which means operations have a nice worst-case logarithmic performance.

Most of the array nodes would only be sparsely populated. To not waste space, the arrays are actually used in conjunction with a bitmap of length 2^k that encodes which

¹A more detailed breakdown of this benchmark can be found in Section 3.4.

positions in the array are actually filled. If a bit is set in the bitmap, then the (logical) array contains an element at the corresponding index. The actual array only has a size equal to the bit count of the bitmap, and after obtaining a (logical) array index i in the manner described above, it has to be converted to an index into the sparse array via the formula $\#((i - 1) \wedge bmp)$, where $\#$ is a function that counts the number of bits and bmp is the array's bitmap. To ensure that the bitmap can be efficiently represented, k is usually chosen so that 2^k equals the size of the native machine word, e.g. on 64-bit systems $k = 6$.

The *concurrent* trie extends the hash trie by adding *indirection nodes* above every array node. An indirection node simply points to the array node underneath it. Indirection nodes have the property that they stay in the trie even if the nodes above or below them change. When inserting an element into the trie, instead of directly modifying an array node, an updated copy of the array node is created and an atomic compare-and-swap operation on the indirection node is used to switch out the old array node for the new one. If the compare-and-swap operation fails, meaning another thread has already modified the array while we were not looking, the operation is retried from the beginning. This simple scheme, where indirection nodes act as barriers for concurrent modification, ensures that there are no lost updates or race conditions of any kind, while keeping all operations completely lock-free. A more thorough discussion, including proofs of linearizability and lock-freedom, can be found in the paper by Prokopec et al. (2011). A Haskell implementation of the concurrent trie, as a mutable data structure in IO, is also available (Schröder 2014).

The *transactional* trie is an attempt to lift the concurrent trie into an STM context. The idea is to use the lock-freedom of the concurrent trie to make a non-contentious data structure for STM. This is not entirely straightforward, as there is a natural tension between the atomic compare-and-swap operations of the concurrent trie, which are pessimistic and require execution inside the IO monad, and optimistic transactions as implemented by STM. While it is possible to simulate compare-and-swap using TVars and `retry`,² this would entangle the indirection nodes with the rest of the transaction, which is exactly the opposite of what we want. To keep the non-blocking nature of the concurrent trie, the indirection nodes need to be kept independent of the transaction as a whole, which should only hinge on the actual values stored in the trie's leaves. If two transactions were to cross paths at some indirection node, but otherwise concern independent elements of the trie, then neither transaction should have to retry or block. Side-effecting compare-and-swap operations that run within but independently of

²Like this, for example:

```
stmCAS :: TVar a -> a -> a -> STM ()
stmCAS var old new = do
  cur <- readTVar var
  if (cur == old)
    then writeTVar var new
    else retry
```

a transaction are the only way to achieve this. Alas, the type system, with good reason, will not just allow us to mix IO and STM actions, so we have to circumvent it from time to time using `unsafeIOToSTM`. We will need to justify every single use of `unsafeIOToSTM` and ensure it does not lead to violations of correctness. Still, bypassing the type system is usually a bad sign, and indeed we will see that correctness can only be preserved at the cost of memory efficiency, at least in an STM implementation without finalizers.

3.2 Implementation

The version of the transactional trie discussed in this chapter is available on Hackage at <http://hackage.haskell.org/package/ttrie-0.1.2>. The full source code can also be found at <http://github.com/mcschroeder/ttrie>.

The module `Control.Concurrent.STM.Map`³ exports the transactional trie under the following interface:

```
data Map k v
empty :: STM (Map k v)
insert :: (Eq k, Hashable k) => k -> v -> Map k v -> STM ()
lookup :: (Eq k, Hashable k) => k -> Map k v -> STM (Maybe v)
delete :: (Eq k, Hashable k) => k -> Map k v -> STM ()
```

Now let us implement it. As always, we begin with some types:⁴

```
newtype Map k v = Map (INode k v)
type INode k v = IRef (Node k v)
data Node k v = Array !(SparseArray (Branch k v))
                | List  ![Leaf k v]
data Branch k v = I !(INode k v)
                | L  !(Leaf k v)
data Leaf k v = Leaf !k !(TVar (Maybe v))
```

The transactional trie largely follows the construction of a concurrent trie:

- The `INode` is the indirection node described in the previous section and is simply an `IRef`, which is a mutable variable in IO. To read and write `IRefs` atomically, we will use some functions and types from the `atomic-primops` package (Newton 2014):

```
data Ticket a
readForCAS :: IRef a -> Ticket a
```

³The name of the trie's public data type is `Map`, instead of, say, `TTrie`. The more general name is in keeping with other container libraries and serves to decouple the interface from the specific implementation based on concurrent tries.

⁴The `!` operator is a strictness annotation.


```

peekTicket :: Ticket a → IO a
casIORef :: IORef a → Ticket a → a → IO (Bool, Ticket a)

```

The idea of the `Ticket` type is to encapsulate proof that a thread has observed a specific value of an `IORef`. Due to compiler optimizations, it would not be safe to just use pointer equality to compare values directly.

- A `Node` is either an `Array` of `Branches` or a `List` of `Leafs`. The `List` is used in case of hash collisions. A couple of convenience functions help us manipulate such collision lists:

```

listLookup :: Eq k ⇒ k → [Leaf k v] → Maybe (TVar (Maybe v))
listDelete :: Eq k ⇒ k → [Leaf k v] → [Leaf k v]

```

Their implementations are entirely standard.

The `Array` is actually a `SparseArray`, which abstracts away all the bit-fiddling necessary for navigating the bit-mapped arrays underlying a hash array mapped trie. Its interface is largely self-explanatory:

```

data SparseArray a
emptyArray :: SparseArray a
mkSingleton :: Level → Hash → a → SparseArray a
mkPair :: Level → Hash → a → a → Maybe (SparseArray a)
arrayLookup :: Level → Hash → SparseArray a → Maybe a
arrayInsert :: Level → Hash → a → SparseArray a → SparseArray a
arrayUpdate :: Level → Hash → a → SparseArray a → SparseArray a

```

I will not go into the implementation of `SparseArray`. It is fairly low-level and can be found in the internal `Data.SparseArray` module of the `ttrie` package.

Some additional functions are used to manipulate `Hashes` and `Levels`. Again, they are self-explanatory:

```

type Hash = Word
hash :: Hashable a ⇒ a → Hash

type Level = Int
down :: Level → Level
up :: Level → Level
lastLevel :: Level

```

- A `Branch` either adds another level to the trie by being an `INode` or it is simply a single `Leaf`.

The one big difference to a concurrent trie lies in the definition of the `Leaf`. Basically, a `Leaf` is a key-value mapping. It stores a key `k` and a value `v`. But the way it stores `v` determines how the trie behaves in a transactional context. Let us build it step by step:

1. Imagine if `Leaf` were defined exactly like in a concurrent trie:

```
data Leaf k v = Leaf !k v
```

Then an atomic compare-and-swap on an `Inode` to insert a new `Leaf` would obviously not be safe during an STM transaction: other transactions could see the new value `v` before our transaction commits; and they could replace `v` by inserting a new `Leaf` for the same key, resulting in our insert being lost.

2. We can eliminate lost inserts by wrapping the value in a `TVar`:

```
data Leaf k v = Leaf !k !(TVar v)
```

Now, instead of replacing the whole `Leaf` to update `v`, we can use `writeTVar` to only modify the value part of the `Leaf`. If two transactions try to update the same `Leaf`, then STM will detect the conflict and one of the transactions would have to retry.

Of course, if there does not yet exist a `Leaf` for a specific key, then a new `Leaf` will still have to be inserted with a compare-and-swap. In this case it is again possible for other transactions to read the `TVar` immediately after the swap, even though our transaction has not yet committed and may still abort. This can happen without conflict because the new `Leaf` contains a newly allocated `TVar` and allocation effects are allowed to escape transactions by design (see Section 2.4). Reading a newly allocated `TVar` will never cause a conflict.

3. To ensure proper isolation, the actual type of `Leaf` looks like this:

```
data Leaf k v = Leaf !k !(TVar (Maybe v))
```

By adding the `Maybe`, we can allocate new `TVars` with `Nothing` in them. A transaction can then insert a new `Leaf` containing `Nothing` using the compare-and-swap operation. Other threads will still be able to read the new `TVar` immediately after the compare-and-swap, but all they will get is `Nothing`. The transaction, meanwhile, can simply `writeTVar (Just v)` to safely insert the actual value into the `Leaf`'s `TVar`. If another transaction also writes to the `TVar` and commits before us, then we have a legitimate conflict on the value level, and our transaction will simply retry.

Now that we have the types that make up the trie's internal structure, we can implement its operations. We begin with the function to create an empty trie:

```
empty :: STM (Map k v)
empty = unsafeOToSTM $ Map <$> newIORef (Array emptyArray)
```

It contains no surprises, although it has the first use of `unsafeOToSTM`, which in this case is clearly harmless.

For the rest of the operations, let us assume we have a function

$\text{getTVar} :: (\text{Eq } k, \text{Hashable } k) \Rightarrow k \rightarrow \text{Map } k \ v \rightarrow \text{STM } (\text{TVar } (\text{Maybe } v))$

that either returns the `TVar` stored in the `Leaf` for a given key, or allocates a new `TVar` for that key and inserts it appropriately into the trie. The `TVar` returned by `getTVar k m` will always either contain `Just v`, where v is the value associated with the key k in the trie m , or `Nothing`, if k is not actually present in m . Additionally, `getTVar` obeys the following invariants:

Invariant 1: $\text{getTVar } k_1 \ m \equiv \text{getTVar } k_2 \ m \iff k_1 \equiv k_2$

Invariant 2: `getTVar` itself does not read from nor write to any `TVars`.

Now we can define the trie's operations as follows:

```
insert k v m = do var ← getTVar k m
                  writeTVar var (Just v)
lookup k m = do var ← getTVar k m
                readTVar var
delete k m = do var ← getTVar k m
               writeTVar var Nothing
```

The nice thing about defining the operations this way, is that correctness and non-contentiousness follow directly from the invariants of `getTVar`. The first invariant ensures correctness. If we get the same `TVar` every time we call `getTVar` with the same key, and if that `TVar` is unique to that key, then STM will take care of the rest. And if, by the second invariant, `getTVar` does not touch any transactional variables, then the only way one of the operations can cause a conflict is if it actually operates at the same time on the same `TVar` as another transaction. Unnecessary contention is therefore not possible.

All that is left to do is implementing `getTVar`. Essentially, `getTVar` is a combination of the `insert` and `lookup` functions of the concurrent trie, just lifted into STM. It tries to look up the `TVar` associated with a given key, and if that does not exist, allocates and inserts a new `TVar` for that key. When inserting a new `TVar`, the structure of the trie has to be changed to accommodate the new element.

Let us look at the code:

```
getTVar k (Map root) = go root 0
  where
    h = hash k
```

The actual work is done by the recursive helper function `go`. It begins at level 0 by looking into the `root` indirection node. Note that throughout the iterations of `go`, the hash h of the key is only computed once.

```
go inode level = do
  ticket ← unsafeIOToSTM $ readForCAS inode
```

```

case peekTicket ticket of
  Array a → case arrayLookup level h a of
    Just (I inode2) → go inode2 (down level)
    Just (L leaf2@(Leaf k2 var))
      | k ≡ k2 → return var
      | otherwise → cas inode ticket (growTrie level a (hash k2) leaf2)
    Nothing → cas inode ticket (insertLeaf level a)
  List xs → case listLookup k xs of
    Just var → return var
    Nothing → cas inode ticket (return ∘ List ∘ (:xs))

```

The use of `unsafeIOToSTM` here is clearly safe—all we are doing is reading the value of the indirection node. This does not have any side effects, so it does not matter if the transaction aborts prematurely. If the transaction retries, the indirection node is just read again—possibly resulting in a different value. It is also possible that the value of the indirection node changes during the runtime of the rest of the function—but that is precisely why we obtain a `Ticket`.

Depending on the contents of the indirection node, we either go deeper into the trie with a recursive call of `go`; return the `TVar` associated with the key; or insert a new `TVar` by using the `cas` function to swap out the old contents of the indirection node with an updated version that somehow contains the new `TVar`.

The `cas` function is also part of the **where** clause of `getTVar`:

```

cas inode ticket f = do
  var ← newTVar Nothing
  node ← f (Leaf k var)
  (ok, _) ← unsafeIOToSTM $ casIORef inode ticket node
  if ok then return var
  else go root 0

```

It implements a transactionally safe compare-and-swap procedure:

1. Allocate a new `TVar` containing `Nothing`.
2. Use the given function *f* to produce a *node* containing a `Leaf` with this `TVar`.
3. Use `casIORef` to compare-and-swap the old contents of the *inode* with the new *node*.
4. If the compare-and-swap was successful, the new *node* is immediately visible to all other threads. Return the `TVar` to the caller, who is now free to use `writeTVar` to fill in the final value.
5. If the compare-and-swap failed, because some other thread has changed the *inode* since the time we first read it, restart the operation—not with the `STM retry`, which would restart the whole transaction, but simply by calling `go root 0` again.

All that is remaining now are the functions given for f in the code of `go`. Given a new `Leaf`, they are supposed to return a `Node` that somehow contains this new `Leaf`. In the case of the overflow list, this is just a trivial anonymous function that prepends the leaf into the `List` node. The `insertLeaf` function does pretty much the same, except for `Array` nodes:

```
insertLeaf level a leaf = do
  let a' = arrayInsert level h (L leaf) a
  return (Array a')
```

In case of a key collision, things are a slightly more involved. The `growTrie` function puts the colliding leaves into a new level of the trie, where they hopefully will not collide anymore:

```
growTrie level a h2 leaf2 leaf1 = do
  inode2 ← unsafeIOToSTM $ combineLeaves (down level) h leaf1 h2 leaf2
  let a' = arrayUpdate level h (I inode2) a
  return (Array a')
combineLeaves level h1 leaf1 h2 leaf2
  | level ≥ lastLevel = newIORef (List [leaf1, leaf2])
  | otherwise =
    case mkPair level h (L leaf1) h2 (L leaf2) of
      Just pair → newIORef (Array pair)
      Nothing → do
        inode ← combineLeaves (down level) h1 leaf1 h2 leaf2
        let a = mkSingleton level h (I inode)
        newIORef (Array a)
```

The use of `casIORef` here is once again harmless, as `combineLeaves` only uses `IO` to allocate new `IORefs`. The `mkPair` function for making a two-element `SparseArray` returns a `Maybe`, because it is possible that on a given level of the trie the two keys hash to the same array index and so the leaves cannot both be put into a single array. In that case, another new indirection node has to be introduced into the trie and the procedure repeated. If at some point the last level has been reached, the leaves just go into an overflow `List` node.

3.3 Memory efficiency

While the transactional trie successfully carries over the lock-freedom of the concurrent trie and keeps the asymptotic performance of its operations, it does have to make a couple of concessions regarding memory efficiency.

The first concession is that when looking up any key for the first time, the `lookup` operation will actually grow the trie. This is a direct consequence of using `getTVar` to implement the trie's basic operations. If `getTVar` does not find the `Leaf` for a given key,

it allocates a new one and inserts it. One might wonder if it is possible to implement a lookup function that does not rely on `getTVar`. The following attempt is pretty straightforward and appears to be correct at first glance—although you might already guess from its name that something is not quite right:

```
phantomLookup :: (Eq k, Hashable k) => k -> Map k v -> STM (Maybe v)
phantomLookup k (Map root) = go root 0
  where
    h = hash k
    go inode level = do
      node <- unsafeIOToSTM $ readIORef inode
      case node of
        Array a -> case arrayLookup level h a of
          Just (I inode2) -> go inode2 (down level)
          Just (L (Leaf k2 var))
            | k ≡ k2    -> readTVar var
            | otherwise -> return Nothing
          Nothing       -> return Nothing
        List xs -> case listLookup k xs of
          Just var -> readTVar var
          Nothing  -> return Nothing
```

The problem with this simple implementation is that under certain circumstances it allows for *phantom reads*. Consider the following pair of functions:

```
f = atomically $ do v1 <- phantomLookup k
                   v2 <- phantomLookup k
                   return (v1 ≡ v2)

g = atomically (insert k 23)
```

Due to STM’s isolation guarantees, one would reasonably expect that *f* always returns `True`. However, sometimes *f* will return `False` when *g* is run between the two `phantomLookup`s in *f*. How is this possible? If you start out with an empty trie, then the first `phantomLookup` in *f* obviously returns `Nothing`. And it does so without touching any TVars, because there is no TVar for *k* at this point. Only when running *g* for the first time, will a TVar for *k* be created. The transaction inside *f* will now happily read from this TVar during the second `phantomLookup` and will not detect any inconsistencies, because this is the first time it has seen the TVar. This problem does not only occur on an empty trie, but any time we look up a key that has not previously been inserted. The only remedy is to ensure that there is always a TVar for every key, even if it is filled with `Nothing`, which is exactly what the implementation of `lookup` using `getTVar` does.

Granted, it seems as if these kinds of phantom lookups might not occur regularly in practice, and even if they did, they would probably cause no great harm. The overhead of always allocating a `Leaf` for every key that is ever looked up, on the other hand,

seems much more troublesome. However, `phantomLookup` exhibits exactly the kind of seldom-occurring unexpected behavior that results in bugs that are incredibly hard to find. And having a `lookup` function that grows the trie is really only an issue in two cases:

1. when we expect the keys we look up to not be present a significant amount of the time; then a transactional trie is probably really not the right data structure. Although if one were to use `phantomLookup` instead of `lookup`, and if in this particular scenario phantom lookups are actually acceptable, then using a transactional trie could still be feasible.
2. when a malicious actor purposefully wants to increase memory consumption, i.e. a classic denial-of-service attack; then one can again counteract this by using `phantomLookup`, limited to those places that are susceptible to attack. For example, a login routine in a web application could first use `phantomLookup` to check if the user actually exists, before continuing with the transaction. Here the phantom lookup does not matter, because if the user does not exist the transaction is aborted anyway.

Thus, it makes sense to have the behavior of `lookup` be the default and provide `phantomLookup` for those select scenarios where it is actually an improvement.

The other trade-off the trie has to make regarding memory efficiency, is that the `delete` operation does not actually remove `Leafs` or compact the trie again. It merely fills a `Leaf`'s `TVar` with `Nothing`. This frees up the values associated with the keys, which is the major part of a trie's memory consumption, but it does not delete the keys or compress the structure that has emerged in the trie, which might now be suboptimal given the trie's current utilization.

Again, for the common use case, this might not be a problem. Very often, we do not want to actually delete certain data, but merely mark it as deleted; or maybe delete the data, but mark the associated keys as having been previously in use in order to prevent reusing them. Think of unique user IDs, for example. In such a scenario, the overhead of the trie not actually deleting `Leafs` disappears. Still, there are of course cases where we do want the trie to always be as compact a representation of its data as possible, and there is in fact a way to achieve this: by using finalizers. What prevents us from just removing `Leafs` from the trie during a transaction is that transactions might get restarted or aborted. If we use the normal `delete` function during the transaction, which essentially just marks a key as deleted (but by accessing the `TVar` ensures that there are no conflicts with other transactions), we can then use an `unsafeDelete` function inside the transaction's finalizer to really remove the `Leaf`:

```
atomicallyWithIO (delete k m) (\_ → unsafeDelete k m)
```

The finalizer ensures the atomicity of the otherwise unsafe operation.

To implement `unsafeDelete` properly and preserve lock-freedom, we have to slightly alter the `Node` type by adding an additional kind of node: a `Tomb` node.

```

data Node k v = Array !(SparseArray (Branch k v))
                  | List  ![Leaf k v]
                  | Tomb  !(Leaf k v)

```

A Tomb node holds a single key. It comes into existence when `unsafeDelete` would result in an `Array` node with only a single `Leaf` beneath it. A Tomb node is the last value assigned to an `INode`. If any operation encounters an `INode` that points to a Tomb node, it cannot modify the `INode` but rather must help compress the trie by merging the tombled leaf into the parent `INode` before retrying the operation. The exact details of `unsafeDelete` and the accompanying cleanup and compression procedures are very tricky, but not especially interesting. They are described in great detail by Prokopec, Bronson, et al. (2012) and are included in the final implementation of the transactional trie.

3.4 Evaluation

I empirically evaluated the transactional trie against similar data structures, measuring contention, runtime performance and memory allocation. The benchmarks were run on an Amazon EC2 C3 extra-large instance with Intel Xeon E5-2680 v2 (Ivy Bridge) processors and a total of 16 physical cores. Under comparison were three hashing-based container types: a transactional trie; a `HashMap` from the `unordered-containers` library (Tibell and Yang 2014), wrapped inside a `TVar`; and the STM-specialized hash array mapped trie from the `stm-containers` library (Volkov 2014b).

Each benchmark consists of a number of random STM transactions. The benchmarks differ in the size and composition of these transactions. Each benchmark is run on every container type, using the same random Text strings as keys each time. The benchmarks are run multiple times, using an increasing number of threads. The transactions are split evenly over the number of threads in use. The time it takes to complete all transactions for a particular container is measured using the `criterion` and `criterion-plus` libraries (O’Sullivan 2014; Volkov 2014a), which calculate the mean execution time over many iterations. To measure contention, the transactions are run again using the `stm-stats` library (Leuschner, Wehr, and Breitner 2011) to count how often the STM runtime system has to restart transactions due to conflicts. Finally, the transactions are run once more to measure the total amount of allocated memory, using GHCs built-in facilities for collecting memory usage statistics. All benchmarks were compiled using GHC 7.8.3. For more details about test data generation and the exact benchmark setup, see the `ttrie` source distribution.

Single-operation transactions. The first four benchmarks (Figure 3.1) each perform 200 000 transactions, where each transaction is just a single operation: insert, update, lookup or delete. The insertion benchmark starts out with an empty container, while all other benchmarks operate on containers prefilled with 200 000 entries. An update operation is simply an insert where the key is already present in the container.

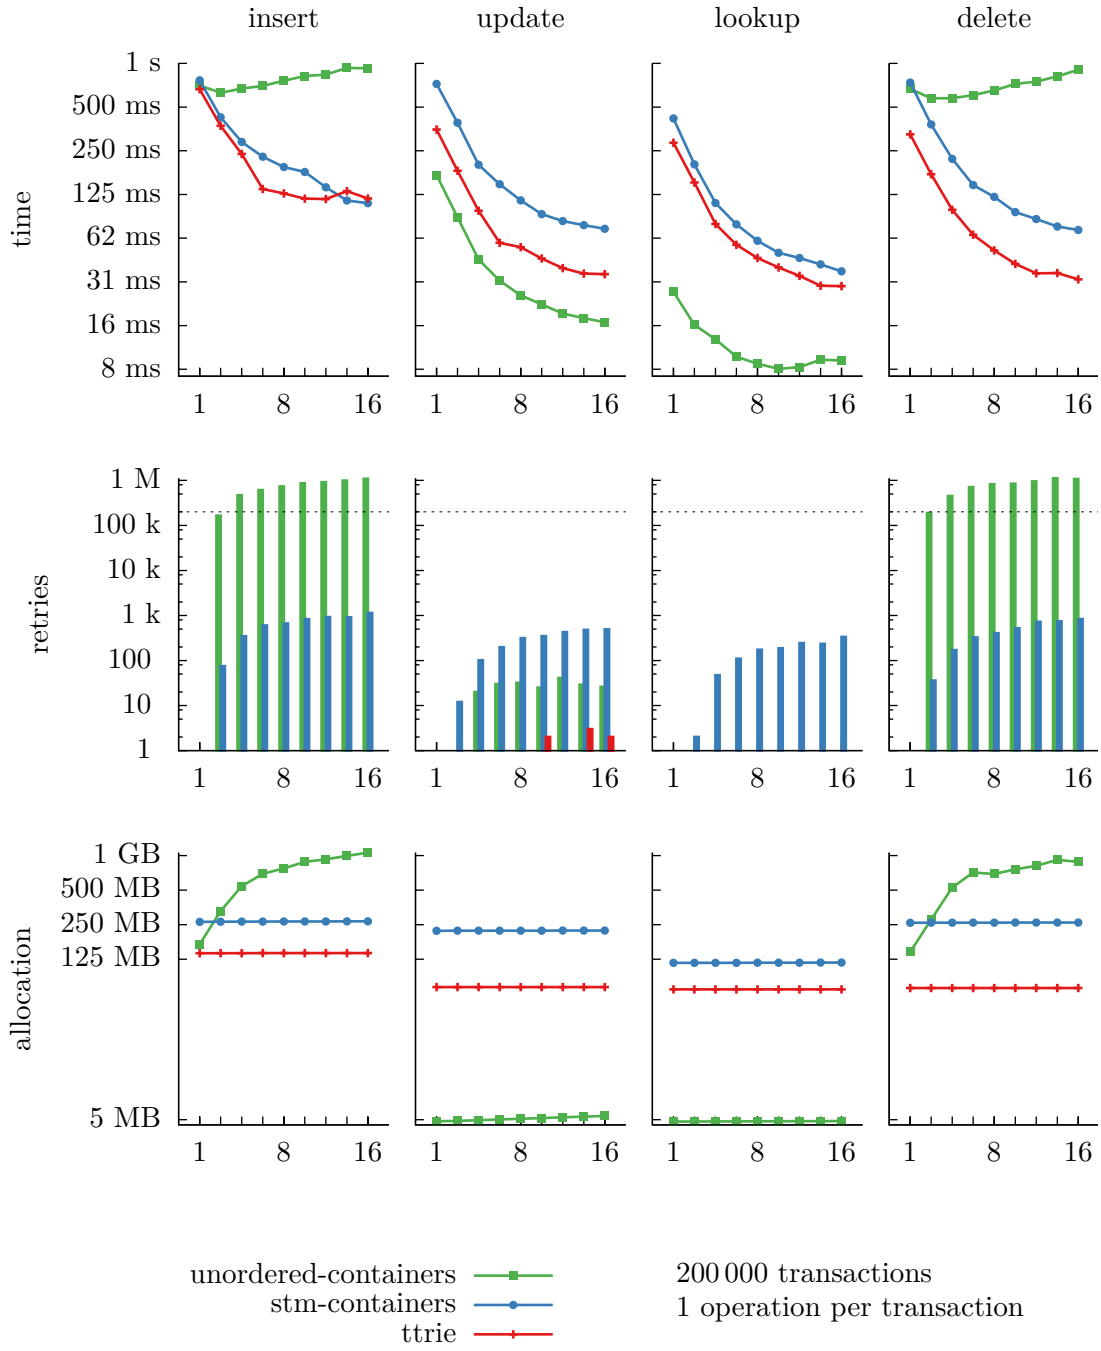


Figure 3.1: Single-operation transactions

The simple TVar-wrapped `HashMap` from `unordered-containers` performs exceptionally well for lookups. This is to be expected: the transactional overhead for reading a single TVar is practically non-existent and the STM runtime system can perform read-only transactions completely lock-free. The greater complexity of `ttrie` and `stm-containers` results in a greater overhead; they perform up to 10 times slower, although they scale pretty well with the number of threads; `ttrie` is about 30% faster than `stm-containers`. Curiously, `stm-containers` exhibits some contention for this read-only operation.

Updates are similarly well-suited to the `HashMap`. Since the keys are already present in the map, there are no structural changes necessary. The transactions are fast enough so that even though there is a small amount of contention — as all transactions have to go through a single TVar — the number of retries stays low enough to not matter. The `HashMap` is roughly twice as fast as the `ttrie`, which is roughly twice as fast in this scenario as the `stm-containers` map.

The story looks entirely different for the insert and delete operations. Here, the TVar-wrapped `HashMap` does not scale at all. On the highest number of threads, it is an order of magnitude slower than the next best contender. The reason is obvious: the amount of contention is so high that the number of retries actually exceeds the number of transactions, i.e. every transaction has to retry at least once; additional threads are actually detrimental to performance. With 16 threads, each one performing only 12 500 transactions, the insert and delete benchmarks recorded over 1 million retries for `unordered-containers`. The transactional trie, as expected, exhibits no contention at all. It is not only much faster than `unordered-containers`, it is also twice as fast as `stm-containers` during all delete benchmarks and about half of the insert benchmarks. The total amount of memory allocated by `unordered-containers` also increases dramatically with the number of threads, while `ttrie` and `stm-containers` both use a constant amount of memory irrespective of the level of concurrency, with `ttrie` using somewhat less.

That the transactional trie is faster than `unordered-containers` during delete even on a single thread, where there can be no contention, is explained by the transactional trie not actually compressing itself after removing a value. This makes a deletion in the trie just a special case of an update. In fact, as long as there are no legitimate conflicts on the value level or interactions within a larger transaction, the update, lookup and delete operations of the `ttrie` should always exhibit the same run-time performance.

Mixed transactions. For the second set of benchmarks (Figure 3.2), transactions are no longer just a single operation, but composed of a mix of up to 5 operations. Using transactions of varying sizes and compositions much closer reflects real-world usage. The insert benchmark now consists of 70% inserts, with the remaining 30% evenly distributed among updates, lookups and deletes. Likewise, the update benchmark now consists of 70% updates, the lookup benchmark of 70% lookups and the delete benchmark of 70% deletes. The insert benchmark again starts out with an empty container, while the other benchmarks operate on containers prefilled with 1 000 000 entries.

As expected, legitimate conflicts between transactions are now slightly more common, evidenced by the increased number of retries measured for the transactional trie,

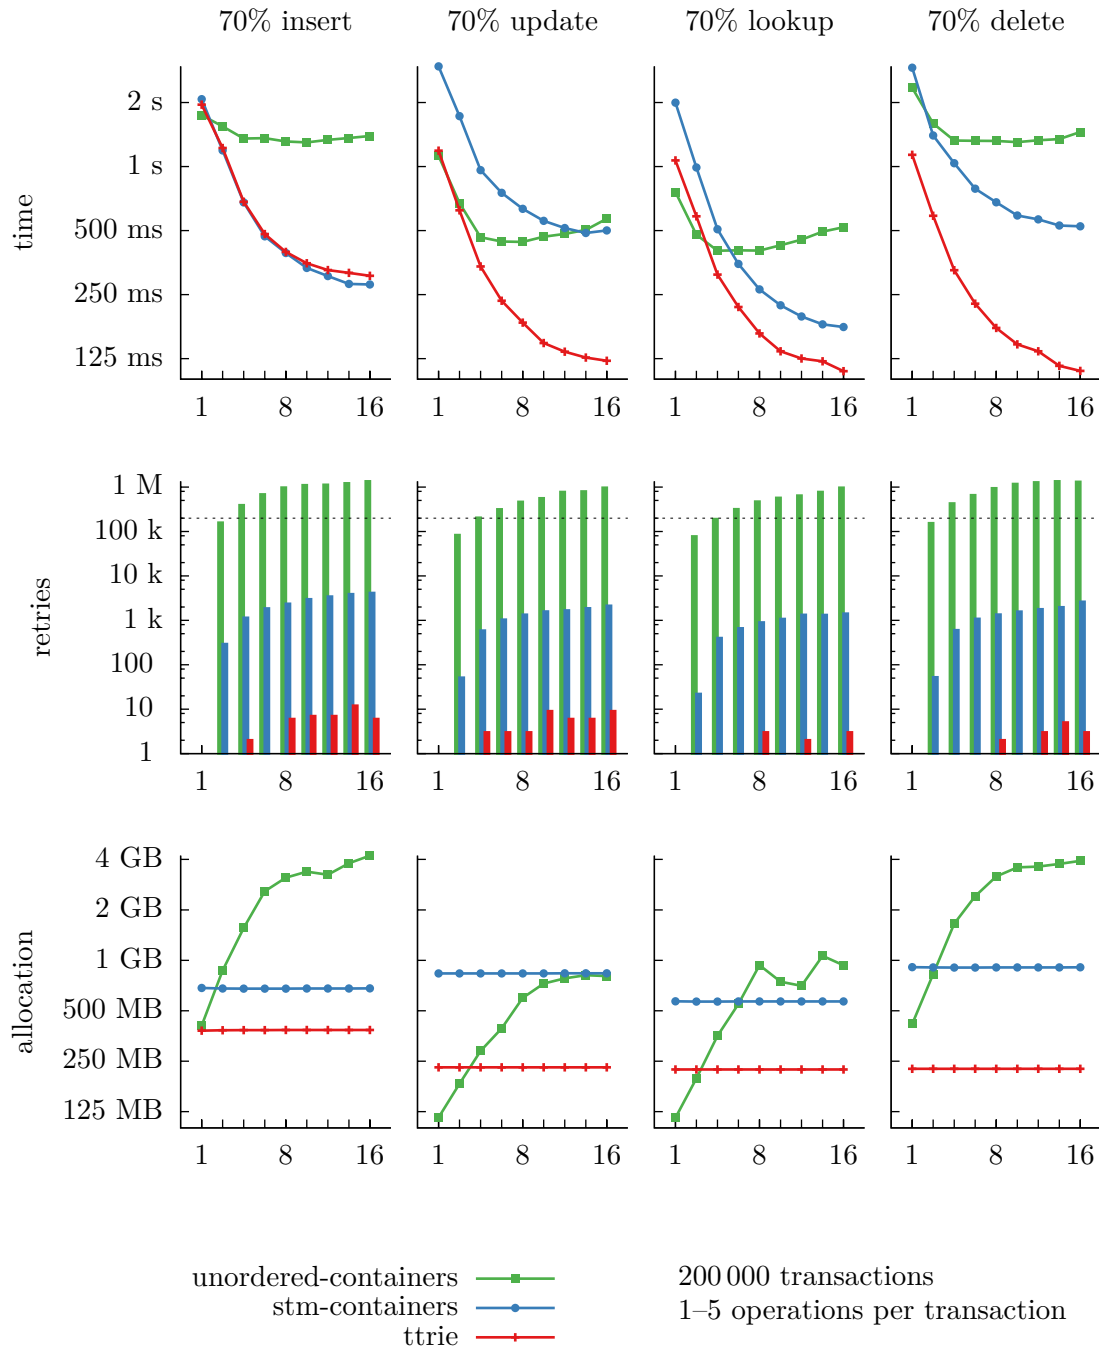


Figure 3.2: Mixed transactions

but they still amount to less than a dozen in the worst case. Yet it should come as no surprise that for `unordered-containers` and `stm-containers` the number of spurious retries vastly overshadows the legitimate conflicts. In the worst case, the TVar-wrapped `HashMap` has to retry more than 1.3 million times for 200 000 transactions to succeed. By mixing in just 30% different operations and using slightly varying transaction sizes, `unordered-containers` exhibits contention even in the predominantly update and lookup scenarios. The run-time performance of `unordered-containers` begins to rapidly degrade at 4 threads, which is when the number of retries first exceeds the number of transactions.

The other results are largely the same as in the first set of benchmarks, just a bit more pronounced. For example, `ttrie` and `stm-containers` now have a virtually identical runtime performance for insert, while the lead `ttrie` had on `stm-containers` during update, lookup and delete has become bigger.

Balanced transactions. The last benchmark (Figure 3.3) consists of a balanced mix of 25% of each operation, on containers prefilled with 1 000 000 entries.

This kind of benchmark plays to the strengths of the transactional trie: here, `ttrie` is 2–4 times faster than `stm-containers`, allocating only a third of the memory; and 1.3–8.6 times faster than `unordered-containers`, allocating almost 10 times less memory.

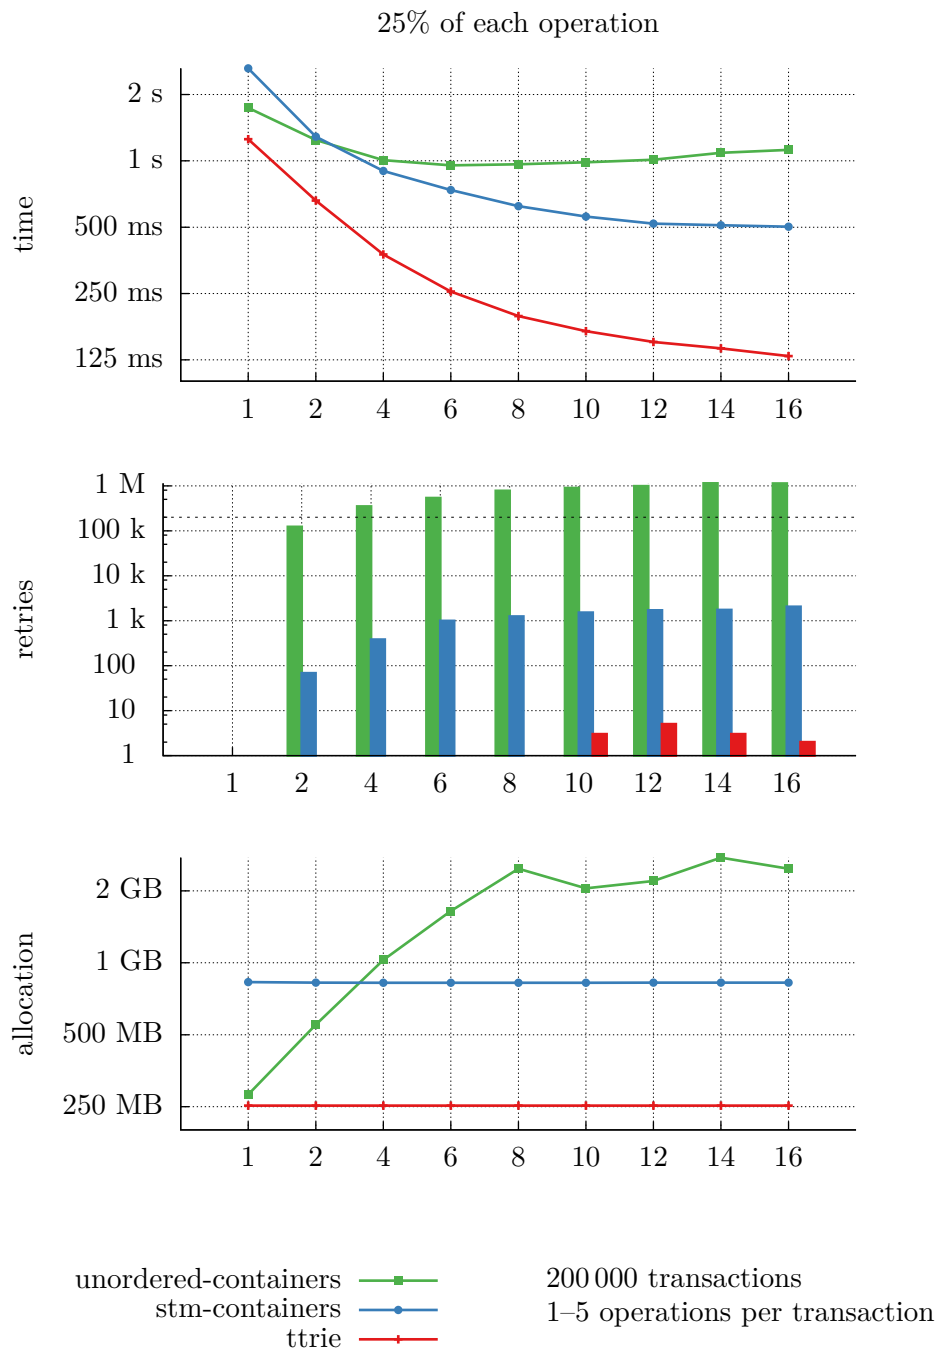


Figure 3.3: Balanced transactions

STM as a database language

As a qualitative demonstration of the effectiveness of both finalizers and transactional tries, I will now use them to build a real application together with a reusable framework for constructing lightweight databases.

This is partly a continuation of my previous work on the `tx` library (Schröder 2013), which was the original motivation for finalizers. The goal of that library was to add a thin persistence layer on top of STM. But without finalizers, this could not be done in a reliable manner, for the reasons laid out in Section 2.1. Now, however, we can do it right.

4.1 Example: a social network

As a running example, we will build a simple social networking site. The full application includes a Haskell web server that exposes a RESTful API and a simple JavaScript client. Here we will focus on the site’s back-end, i.e. its data types and business logic. All code snippets in this chapter are taken from fully working programs. The complete sample code is available at <http://github.com/mcschroeder/social-example>.

Our social network will start out with a modest set of features: users can post messages to their *timelines*; users can *follow* other users; and each user has a personalized *feed*, which interweaves her timeline with the timelines of all the people she follows.

Because we want to quickly arrive at a working prototype, we use STM for concurrency and postpone the question of durability until later. This way we can focus on getting the types right for our business logic, without having to worry about interfacing with an external database just now. This first version of the social network can be found in the `social0` folder of the sample code.

Speaking of types, here they are:

```

data SocialDB = SocialDB
  { users :: Map UserName User
  , posts :: Map PostId Post
  }

type UserName = Text

data User = User
  { name    :: UserName
  , timeline :: TVar [Post]
  , following :: TVar (Set User)
  , followers :: TVar (Set User)
  }

newtype PostId = PostId Word64
  deriving (Eq, Ord, Random, Show, Hashable)

data Post = Post
  { postId :: PostId
  , author :: User
  , time   :: UTCTime
  , body   :: Text
  }

```

The whole of the network is contained in the `SocialDB` type. Users are identified by their names and are represented by the `User` type. Each user keeps a list of the posts on her own timeline and also keeps track of other users that she follows and that are following her. Posts are identified by a globally unique `PostId` and are represented by the `Post` type, which contains the body of the post as well as its author and the time it was created.

Behind the `Map` type of the `users` and `posts` collections lies the transactional trie from Chapter 3. This guarantees that there will be no contention when our millions of future users all post something to our site at the same time.

Since the types contain transactional variables, computations on them must be done in the STM monad. For example, this is how we compute a user's feed:

```

feed :: User → STM [Post]
feed user = do
  myPosts ← readTVar (timeline user)
  others ← Set.toList <$> readTVar (following user)
  otherPosts ← concat <$> mapM (readTVar ∘ timeline) others
  return $ sortBy (flip $ comparing time) (myPosts ++ otherPosts)

```


And this is how users can follow each other:

```
follow :: User → User → STM ()
follow user1 user2 = do
  modifyTVar (following user1) (Set.insert user2)
  modifyTVar (followers user2) (Set.insert user1)
```

By using STM, we can write high-level code and get atomicity for free. Additionally, we can take advantage of some other nice STM features, like composable blocking:

```
waitForFeed :: User → UTCTime → STM [Post]
waitForFeed user lastSeen = do
  posts ← takeWhile isNew <$> feed user
  if null posts then retry else return posts
where
  isNew post = diffUTCTime (time post) lastSeen > 0.1
```

The `retry` operator enables `waitForFeed` to block until there are new posts in a user's feed. Our server uses this function to effortlessly push updates to the client in real-time via HTTP long polling (Loreto et al. 2011).

Creating new posts is also relatively straightforward:

```
createPost :: User → Text → SocialDB → STM Post
createPost author body db = do
  postId ← newUniquePostId db
  time ← unsafeIOToSTM getCurrentTime
  newPost postId author time body db

newUniquePostId :: SocialDB → STM PostId
newUniquePostId db = do
  postId ← unsafeIOToSTM randomIO
  alreadyExists ← Map.member postId (posts db)
  check (¬ alreadyExists)
  return postId

newPost :: PostId → User → Time → Text → SocialDB → STM Post
newPost postId author time body db = do
  let post = Post {..}
  modifyTVar (timeline author) (post:)
  Map.insert postId post (posts db)
  return post
```

In order to generate a random ID for the post and to get the current time, we have to use `unsafeIOToSTM` to interleave I/O actions with the transaction. In these cases, this is

perfectly safe: we are either only reading from the outside world (`getCurrentTime`) or are otherwise producing only harmless side effects (`randomIO`, which updates the internal state of its random number generator). If the transaction aborts, no harm is done. If the transaction retries, we simply get a new time and generate a new random ID, which is just what we want.

We have now seen pretty much the whole back-end of the social network, apart from `createUser`, which is not that different from `createPost`. Using STM, we managed to arrive very quickly at a functional prototype. Of course, since we are using *only* STM, all effects stay purely in memory. If the server is shut down, the data is gone. In the next section, we will use finalizers to change that.

4.2 The TX monad

The idea is to use an STM finalizer to record state-changing operations in a write-ahead log file. We do not serialize the data itself, but rather the operations on the data. We record function calls that can later be replayed to restore the state of the system at the time of the recording. A similar scheme has been implemented by the `acid-state` library (Himmelstrup 2014), which uses a custom lock-based state container instead of building on STM. The fact that `acid-state` is used by Hackage, the official Haskell package repository, demonstrates the practicality of such an approach.

To facilitate the recording of operations, a thin layer on top of STM is introduced: the TX monad. Here is the `follow` function from the previous section, rewritten in TX:

```
follow :: User → User → TX SocialDB ()
follow user1 user2 = do
  record $ Follow (name user1) (name user2)
  liftSTM $ do
    modifyTVar (following user1) (Set.insert user2)
    modifyTVar (followers user2) (Set.insert user1)
```

This demonstrates the two main tasks we do in TX: lifting functions from the underlying STM monad using `liftSTM` and recording database operations using `record`. Note how TX is parameterized by the type of database, in our case `SocialDB`.

An updated version of the social network can be found in the `social1` folder of the sample code. Below is `createPost` rewritten for TX. The `getData` function is used to retrieve the database implicitly carried around by TX:

```
createPost :: User → Text → TX SocialDB Post
createPost author body = do
  db ← getData
  postId ← liftSTM $ newUniquePostId db
  time ← unsafeIOToTX getCurrentTime
  record $ NewPost postId (name author) time body
  liftSTM $ newPost postId author time body db
```

Notice that the TX version of `createPost` is virtually identical to the pure STM version, except for a few lifts and the call to `record`. The existing STM functions `newPost` and `newUniquePostId` could be reused and did not have to be touched at all.

The implementation of TX itself is pretty simple. It is a strict state transformer¹ on top of STM, keeping a list of recorded operations, as well as allowing easy access to the root database type:

```
newtype TX d a = TX (StateT (d, [Operation d]) STM a)
deriving (Functor, Applicative, Monad)
record :: Operation d → TX d ()
record op = TX $ modify $ λ(d, log) → (d, op : log)
getData :: TX d d
getData = TX $ gets fst
liftSTM :: STM a → TX d a
liftSTM = TX ∘ lift
```

TX could be implemented in a more strongly typed fashion, using a combination reader/writer monad to emphasize that the database reference `d` is constant and that the operation log is strictly write-only. Alas, using `StateT` is much more efficient, in addition to being more convenient, as writer transformers do not work in constant space.²

The argument to `record` is an `Operation d`, which is an associated type of the `Database` class:

```
class Database d where
  data Operation d
  replay :: Operation d → TX d ()
```

Associated types (Chakravarty et al. 2005) are defined by concrete instances of the class they are associated with. In our case, the `Database` instance of `SocialDB` defines `Operation SocialDB`, denoting operations that can be replayed in the context of `SocialDB`. The constructors of `Operation SocialDB` represent the functions that recorded them. For example, the `Follow` and `NewPost` constructors represent the `follow` and `newPost` functions:

```
instance Database SocialDB where
  data Operation SocialDB = Follow UserName UserName
                        | NewPost PostId UserName UTCTime Text
                        | ...
```

¹The `transformers` package (Gill and Paterson 2014) provides a variety of monad transformers.

²See <https://mail.haskell.org/pipermail/libraries/2013-March/019528.html>.

```

replay (Follow name1 name2) = do
  user1 ← getUser name1
  user2 ← getUser name2
  user1 'follow' user2
replay (NewPost postId name time body) = do
  db ← getData
  author ← getUser name
  liftSTM $ void $ newPost postId author time body db

```

So the follow function records the Follow type and replaying the Follow type causes the follow function to be called again. Note that TVars and types that contain TVars are not directly serializable. Therefore the Follow type contains only the names of the users. During replay, the full User structure is fetched from the database.

Now, how is a TX computation actually performed and its effects serialized? The TX equivalent to atomically is durably, and this is where finalizers finally come into play:

```

durably :: DatabaseHandle d → TX d a → IO a
durably h (TX m) = atomicallyWithIO action finalizer
  where
    action                = runStateT m (database h, [])
    finalizer (a, (_, ops)) = serialize ops h >> return a

```

The durably function runs the state monad inside TX and performs the resulting STM action using atomicallyWithIO, passing the logged operations to the finalizer, which serializes them and returns the transaction’s final result. DatabaseHandle *d* is an opaque type referring to a specific database and its on-disk representation. The application programmer uses the function

```

openDatabase :: (Database d, SafeCopy (Operation d))
              ⇒ FilePath → d → IO (DatabaseHandle d)

```

to acquire a database handle that can be shared freely among threads. Upon opening a database, its in-memory representation *d* is initialized by replaying the stored operations from the underlying log file.

Note the SafeCopy constraint on Operation *d*. It comes from the widely-used [safecopy](#) library (Himmelstrup and Lessa 2014), which is a serialization library with support for version control. By leveraging [safecopy](#), TX supports schema migration right out of the box. The ability to extend and change data types and business logic while keeping compatibility with old database files is very important for any database system, and maybe especially so for something as lightweight as TX, which naturally lends itself to rapid prototyping. For a practical demonstration of this, look in the `social2` directory of the sample code. It contains a version of the social network that has some additional features, such as allowing users to post messages directly to the timelines of other users and adding the ability to “like” posts. Implementing these features necessitated non-trivial changes to the core types, yet switching from an older version of the serialized database is seamlessly possible.

4.3 Caveats

Due to the simplicity of the design and the nature of the underlying STM abstraction, there are some drawbacks and limitations. They are not fundamental, but they indicate that in order to achieve higher scalability (in terms of application complexity), a more sophisticated interface might be needed.

Using exceptions to handle failure. In some of the code above we have used a function `getUser` to retrieve a user from the database by name. But what happens if there is no user for the given name? Usually, we want to make the possibility of failure explicit in the types and so would expect `getUser` to have a type like

```
getUser :: UserName → TX SocialDB (Maybe User)
```

However, consider the following scenario:

```
do eve ← newUser "Eve"
   adam ← getUser "Adam"
   case adam of
     Just adam → eve 'follow' adam
     Nothing   → return ()
```

If `getUser` returns `Nothing`, then even though the block returns without further action, *eve* is still created when the transaction ultimately commits. Of course this particular example is somewhat contrived, as we could easily rearrange the functions to postpone the creation of *eve* until after we have established that *adam* exists. But you can imagine that this may not always be possible and does not scale well.

The real problem is that there is no way to explicitly abort an STM transaction other than throwing an exception. This is why `getUser` is actually implemented like this:

```
getUser :: UserName → TX SocialDB User
getUser name = do
  db ← getData
  liftSTM $ do
    user ← Map.lookup name (users db)
    case user of
      Just user → return user
      Nothing   → throwSTM (UserNotFound name)
```

In the majority of uses, the desired outcome of a function like `getUser` in case the user is not found is to abort the transaction. There are certainly ways to avoid spooky exceptions as much as possible, e.g. by adding an exception monad transformer layer to TX (something like `Control.Monad.Trans.Except` from [transformers](#)). But in the end a real exception would need to be thrown regardless, in order to signal to the STM runtime that the transaction needs to be aborted. I have favored the simpler design and it seems to work well in practice, as indicated by our example application.

Double replay. The `record` function gives us great flexibility in how and when to record operations. But with great flexibility comes great responsibility: not only does the programmer have to remember to actually record a specific operation, she must be especially careful when composing functions that record something. Consider the following situation:

$$\begin{aligned} f &= \text{record } F \gg g \\ g &= \text{record } G \end{aligned}$$

Running f first records F and then calls g and so also records G . Replaying this recording will first replay F , which runs f which calls g , and then it will replay G , which calls g again. We have now called g twice, even though it was only executed once during the original run! This is unfortunate, because it means that to avoid such situations, the user has to keep the call graph of recordable functions in mind, and there is no help from the compiler.

This makes one wonder: why do we record on this semi-high level—that obviously is not high enough to save us from such basic yet easily overlooked mistakes—instead of just directly recording every primitive `writeTVar` operation? We may not be able to serialize a `TVar` itself, but we can certainly serialize its contents. Could we not have a function

$$\text{durableWriteTVar} :: \text{Serializable } a \Rightarrow \text{TVar } a \rightarrow a \rightarrow \text{TX } d ()$$

that automatically logs updates at the lowest level, removing the need for explicit calls to `record`? Alas, this is not possible: during replay, there is no way to automatically re-associate the recorded value with the `TVar` it originally belonged to. We always need some kind of context to find the one place in our data structure that has the correct `TVar` to put the value back into. By recording higher-level operations, this context is automatically included.

Conclusions & Perspectives

I have presented two new approaches of combining transactions with side-effects, with the goal of adding *durability* and eliminating *contention*:

- **STM Finalizers** are a global, top-down approach of adding arbitrary I/O during transactional commit. I have given a precise semantics of this extension to Haskell’s STM system, as well as an implementation in the Glasgow Haskell Compiler. I have given many examples of the usefulness of finalizers, and demonstrated how they make it practical and easy to build a lightweight database application on top of STM.
- The **transactional trie** is a specialized data structure that uses local side-effects to eliminate contention inside STM. I have discussed its implementation in detail and reasoned about its safety and related trade-offs. I have shown the superior performance of the transactional trie relative to similar data structures.

5.1 Related Work

Twilight STM (Bieniusa 2011a) is “a transactional memory system [that] augments transactions with non-reversible operations and allows introspection and modification of a transaction’s state.” It has a library-level implementation in Haskell, available in the `twilight-stm` package (Bieniusa 2011b), which provides a custom STM monad. Twilight STM distinguishes between three transactional phases: the *atomic* phase is the same as in standard STM; in the *twilight* phase, inconsistencies between the transactional state and the global state can be detected and fixed, or otherwise ignored; and once in the *safe* phase, a transaction is guaranteed to commit, unless explicitly retried, but the transaction’s write set can still be modified and irrevocable I/O actions can be performed.

The ability to check for and ignore transactional conflicts during the twilight phase enables the implementation of contention free data structures similar to the transactional trie. For example, a lookup operation on a singly-linked list could safely ignore any inconsistencies arising due to concurrent inserts, thereby increasing performance.

The safe phase of Twilight STM is similar to finalizers, and somewhat more powerful: in addition to safely performing I/O actions, it also allows limited reading and writing of transactional variables. Unlike finalizers, however, Twilight STM does not support nesting of transactions inside its safe phase.

Since `twilight-stm` is a standalone library, without GHC runtime support, its performance is drastically reduced compared to native STM. It is also not compatible with existing STM code.

The `AdvSTM` monad from the `stm-io-hooks` package was already mentioned in Section 2.3. In addition to the design differences described there (it is bottom-up rather than top-down; it allows for higher composability, even though this may be problematic in conjunction with I/O), it also has the drawback of being incompatible with existing STM code.

There have been a number of other STM implementations trying to reconcile transactions and side-effects, but all in the context of weakly typed, non-functional languages:

Welc, Saha, and Adl-Tabatabai (2008) describe a transactional memory system for Java that allows transitioning into an irrevocable state in which a transaction will no longer roll back. Any subsequent actions, like I/O effects, will never be revoked or repeated. They use a technique called single-owner read locks and allow only one irrevocable transaction to run at a time.

Harris (2005) extends a Java STM system with external actions, which can perform operations directly on the heap of a given context. When an atomic block finishes, it promotes heap updates from nested contexts up to its parent context.

Spear, Michael, and Scott (2008) compare several mechanisms for “inevitable” transactions, which cannot abort and of which only one can run at a time.

JudoSTM (Olszewski, Cutler, and Steffan 2007) is a system that uses dynamic binary-rewriting to transform C and C++ applications to support transactional execution. It has a concept of “privileged” transactions, which cannot be rolled back and can be used to make system calls.

The xCall interface for the Intel STM compiler (Volos et al. 2009) enables transactions to make system calls, through a combination of delaying the execution of calls until the transaction commits and undoing the side effects of some immediately executed calls should the transaction abort.

Sonmez et al. (2007) extend Haskell’s STM with an `unreadTVar` operation that removes a transactional variable from the read set of a transaction. A data structure like a linked list normally has a read set that is directly proportional to the length of the list, increasing the potential for false conflicts and aborts. By unreading TVars, one can traverse a list with a small fixed-size read set. This results in much faster execution times,

although it can lead to unsafe situations, especially when composing with functions that are not aware of `unreadTVar`.

Transactional boosting (Herlihy and Koskinen 2008) is a methodology for transforming highly-concurrent linearizable objects into highly-concurrent transactional objects, provided they satisfy certain regularity properties. For an object to be boosted, its methods must be commutative, or otherwise be protected by an abstract lock, and they must have fast inverses, which will be executed if the transaction fails to commit; the data type itself is treated like a black box. Du Bois, Pilla, and Duarte (2014) describe a `boost` function for Haskell, as an extension to a high-level STM implementation (Du Bois 2011).

The transactional trie is not a boosted version of the concurrent trie. The lock-free concurrent trie does not seem particularly amenable to transactional boosting, since to ensure commutativity and preserve isolation, some operations would need to acquire locks and introduce additional bookkeeping. The benefits of reusing an existing concurrent trie implementation via transactional boosting are somewhat reduced by the necessity of explicit locking, which decreases performance and adds significantly to overall complexity, on top of the machinery necessary to enable boosting in the first place.

5.2 Future Work

Finalizers allow transactions to be interactive. In particular, they enable interaction with remote systems. Building a distributed transactional memory system on top of finalizers would be the next logical step, leveraging existing distributed computation frameworks such as Cloud Haskell (Epstein 2011; Coutts, Wu, and Vries 2015). It would be interesting to see if and how such a system could be implemented using just finalizers, and how it would compare to more fundamental approaches like the `DSTM` library (Kupke 2010) or Decent STM (Bieniusa 2011a).

By allowing nesting of transactions, finalizers can truly be arbitrary I/O actions, including those that are themselves composed of STM transactions—like parts of the Cloud Haskell platform. But in the current design, programmer error in the form of circular dependencies between inner and outer transactions is only detected at runtime, even though it should be possible to do this statically. One could write a compiler plugin or, more radically, rewrite the STM API to be more strongly typed, perhaps by using an effect system (Orchard and Petricek 2014). Such a new API would be incompatible with current STM code, but a full redesign of the interface could open up other possibilities.

The TX monad could also benefit from a more strongly typed API, as a lot of bookkeeping is still up to the programmer. Different approaches to designing a persistency layer on top of finalizers should be explored. Additionally, there are some low-hanging fruits in the current design: creating checkpoints of the database state at regular intervals and compressing the on-disk log files would save memory and greatly reduce replay time at startup. There is also the possibility of making the database distributed, with or without relying on a wholly distributed STM system underneath.

The concurrent trie, on which the transactional trie is based, supports an efficient non-blocking snapshot operation (Prokopec, Bronson, et al. 2012). This allows, for example, to fold over the trie while it is concurrently modified, without loss of consistency. It is not clear if it is possible to lift this operation into the transactional setting as well, while keeping the same performance characteristics, but further exploration is needed.

Bibliography

- Bagwell, Phil (2001). *Ideal Hash Trees*. Tech. rep. LAMP-REPORT-2001-001. EPFL (cit. on p. 32).
- Bieniusa, Annette (2011a). “Consistency, Isolation, and Irrevocability in Software Transactional Memory.” PhD thesis. University of Freiburg (cit. on pp. 57, 59).
- (2011b). *twilight-stm: STM library with safe irrevocable I/O and inconsistency repair*. Version 1.2. URL: <http://hackage.haskell.org/package/twilight-stm-1.2> (cit. on p. 57).
- Chakravarty, Manuel M. T., Gabriele Keller, Simon Peyton Jones, and Simon Marlow (2005). “Associated Types with Class.” In: *Proceedings of the 32nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. POPL ’05. Long Beach, California: ACM, pp. 1–13 (cit. on p. 53).
- Coutts, Duncan, Nicolas Wu, and Edsko de Vries (2015). *distributed-process: Cloud Haskell: Erlang-style concurrency in Haskell*. Version 0.5.5. URL: <http://hackage.haskell.org/package/distributed-process-0.5.5> (cit. on p. 59).
- Du Bois, André Rauber (2011). “An Implementation of Composable Memory Transactions in Haskell.” In: *Software Composition*. Ed. by Sven Apel and Ethan Jackson. Vol. 6708. Lecture Notes in Computer Science. Springer Berlin Heidelberg, pp. 34–50 (cit. on p. 59).
- Du Bois, André Rauber, Maurício Lima Pilla, and Rodrigo Duarte (2014). “Transactional Boosting for Haskell.” In: *Programming Languages*. Ed. by Fernando Magno Quintão Pereira. Vol. 8771. Lecture Notes in Computer Science. Springer International Publishing, pp. 145–159 (cit. on p. 59).
- Epstein, Jeffrey (2011). “Functional programming for the data centre.” MA thesis. University of Cambridge (cit. on p. 59).
- Gill, Andy and Ross Paterson (2014). *transformers: Concrete functor and monad transformers*. Version 0.4.0.0. URL: <http://hackage.haskell.org/package/stm-stats-0.4.0.0> (cit. on p. 53).
- Harris, Tim (2005). “Exceptions and side-effects in atomic blocks.” In: *Sci. Comput. Program*. 58.3, pp. 325–343 (cit. on p. 58).
- Harris, Tim, Simon Marlow, Simon Peyton Jones, and Maurice Herlihy (2005). “Composable Memory Transactions.” In: *Proceedings of the Tenth ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*. PPOPP ’05. Chicago, IL, USA: ACM, pp. 48–60 (cit. on pp. 3, 11, 13, 15–17).

- Harris, Tim and Simon Peyton Jones (2006). “Transactional memory with data invariants.” In: *First ACM SIGPLAN Workshop on Languages, Compilers, and Hardware Support for Transactional Computing*. TRANSACT ’06. Ottawa, Canada: ACM (cit. on pp. 3, 11, 23).
- Herlihy, Maurice and Eric Koskinen (2008). “Transactional Boosting: A Methodology for Highly-concurrent Transactional Objects.” In: *Proceedings of the 13th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*. PPOPP ’08. Salt Lake City, UT, USA: ACM, pp. 207–216 (cit. on p. 59).
- Himmelstrup, David (2014). *acid-state: Add ACID guarantees to any serializable Haskell data structure*. Version 0.12.2. URL: <http://hackage.haskell.org/package/acid-state-0.12.2> (cit. on p. 52).
- Himmelstrup, David and Felipe Lessa (2014). *safecopy: Binary serialization with version control*. Version 0.8.3. URL: <http://hackage.haskell.org/package/safecopy-0.8.3> (cit. on p. 54).
- Kupke, Frank (2010). *DSTM: A framework for using STM within distributed systems*. Version 0.1.2. URL: <http://hackage.haskell.org/package/DSTM-0.1.2> (cit. on p. 59).
- Leuschner, David, Stefan Wehr, and Joachim Breitner (2011). *stm-stats: retry statistics for STM transactions*. Version 0.2.0.0. URL: <http://hackage.haskell.org/package/stm-stats-0.2.0.0> (cit. on p. 42).
- Loreto, Salvatore, Peter Saint-Andre, Stefano Salsano, and Greg Wilkins (2011). *Known Issues and Best Practices for the Use of Long Polling and Streaming in Bidirectional HTTP*. RFC 6202 (cit. on p. 51).
- Marlow, Simon (2013a). *async: Run IO operations asynchronously and wait for their results*. Version 2.0.1.5. URL: <http://hackage.haskell.org/package/async-2.0.1.5> (cit. on p. 9).
- (2013b). *Parallel and Concurrent Programming in Haskell*. O’Reilly Media Inc. (cit. on p. 3).
- Newton, Ryan (2014). *atomic-primops: A safe approach to CAS and other atomic ops in Haskell*. Version 0.6. URL: <http://hackage.haskell.org/package/atomic-primops-0.6> (cit. on p. 34).
- Olszewski, Marek, Jeremy Cutler, and J. Gregory Steffan (2007). “JudoSTM: A Dynamic Binary-Rewriting Approach to Software Transactional Memory.” In: *Proceedings of the 16th International Conference on Parallel Architecture and Compilation Techniques*. PACT ’07. Brasov, Romania (cit. on p. 58).
- Orchard, Dominic and Tomas Petricek (2014). “Embedding Effect Systems in Haskell.” In: *Proceedings of the 2014 ACM SIGPLAN Symposium on Haskell*. Haskell ’14. Gothenburg, Sweden: ACM, pp. 13–24 (cit. on p. 59).
- O’Sullivan, Bryan (2014). *criterion: Robust, reliable performance measurement and analysis*. Version 1.0.2.0. URL: <http://hackage.haskell.org/package/criterion-1.0.2.0> (cit. on p. 42).
- Prokopec, Aleksander, Phil Bagwell, and Martin Odersky (2011). *Cache-Aware Lock-Free Concurrent Hash Tries*. Tech. rep. EPFL-REPORT-166908. EPFL (cit. on pp. 32, 33).

- Prokopec, Aleksander, Nathan Grasso Bronson, Phil Bagwell, and Martin Odersky (2012). “Concurrent tries with efficient non-blocking snapshots.” In: *Proceedings of the 17th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*. PPOPP 2012. New Orleans, LA, USA: ACM, pp. 151–160 (cit. on pp. 42, 60).
- Robinson, Peter and Chris Kuklewicz (2012). *stm-io-hooks: STM with IO hooks*. Version 0.7.5. URL: <http://hackage.haskell.org/package/stm-io-hooks-0.7.5> (cit. on p. 10).
- Schröder, Michael (2013). *tx: Persistent transactions on top of STM*. Version 0.1.0.0. URL: <http://hackage.haskell.org/package/tx-0.1.0.0> (cit. on p. 49).
- (2014). *ctrie: Non-blocking concurrent map*. Version 0.1.0.2. URL: <http://hackage.haskell.org/package/ctrie-0.1.0.2> (cit. on p. 33).
- Sonmez, Nehir et al. (2007). “UnreadTVar: Extending Haskell Software Transactional Memory for Performance.” In: *Proceedings of the Eighth Symposium on Trends in Functional Programming*. TFP 2007. New York, NY, USA: Intellect (cit. on p. 58).
- Spear, Michael F., Maged Michael, and Michael L. Scott (2008). “Inevitability Mechanisms for Software Transactional Memory.” In: *3rd ACM SIGPLAN Workshop on Transactional Computing*. TRANSACT ’08 (cit. on p. 58).
- Tibell, Johan and Edward Z. Yang (2014). *unordered-containers: Efficient hashing-based container types*. Version 0.2.5.0. URL: <http://hackage.haskell.org/package/unordered-containers-0.2.5.0> (cit. on p. 42).
- Volkov, Nikita (2014a). *criterion-plus: Enhancement of the criterion benchmarking library*. Version 0.1.3. URL: <http://hackage.haskell.org/package/criterion-plus-0.1.3> (cit. on p. 42).
- (2014b). *stm-containers: Containers for STM*. Version 0.2.3. URL: <http://hackage.haskell.org/package/stm-containers-0.2.3> (cit. on p. 42).
- Volos, Haris et al. (2009). “xCalls: Safe I/O in Memory Transactions.” In: *Proceedings of the 4th ACM European Conference on Computer Systems*. EuroSys ’09. Nuremberg, Germany: ACM, pp. 247–260 (cit. on p. 58).
- Welc, Adam, Bratin Saha, and Ali-Reza Adl-Tabatabai (2008). “Irrevocable Transactions and Their Applications.” In: *Proceedings of the Twentieth Annual Symposium on Parallelism in Algorithms and Architectures*. SPAA ’08. Munich, Germany: ACM, pp. 285–296 (cit. on p. 58).
- Yates, Ryan (2013). *GHC Commentary: Software Transactional Memory (STM)*. URL: <http://ghc.haskell.org/trac/ghc/wiki/Commentary/Rts/STM> (cit. on p. 27).

Publications

Presentations

- “Transactional Memory with Finalizers,” presented at the *32nd Workshop of GI SIG Programmiersprachen und Rechenkonzepte*. Bad Honnef, Germany. May 2015
- “Transactional Tries,” to be presented at the *18th Biennial Workshop on Programmiersprachen und Grundlagen der Programmierung*. KPS 2015. Pörschach am Wörthersee, Austria. October 2015

Tools

- A patched version of GHC that supports STM finalizers.
<http://github.com/mcschroeder/ghc>
- The `ttrie` package: *Contention-free STM hash map*.
<http://hackage.haskell.org/package/ttrie>
- A simple social network, built on STM finalizers and transactional tries.
<http://github.com/mcschroeder/social-example>